

# Word Embedding

Ninghao Liu

University of Georgia

March 20, 2024

# 1. How to Represent A Word?

Treating each word as a token.

- In previous methods, each word is represented as an index. Different terms are mutually independent.
- Another perspective: Each word is a **one-hot vector** in the  $|\mathcal{V}|$  dimensional space.

$$\mathbf{w}^a = \begin{bmatrix} 1 \\ 0 \\ \vdots \\ 0 \end{bmatrix}, \quad \mathbf{w}^{at} = \begin{bmatrix} 0 \\ 1 \\ \vdots \\ 0 \end{bmatrix}, \quad \dots, \quad \mathbf{w}^{zoo} = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ 1 \end{bmatrix}. \quad (1)$$

- The similarity issue: Ideally we would want similar words like “learn” and “study” should be similar in the feature space.

# 1. How to Represent the Meaning of A Word?

Traditional NLP solution: Maintaining synonym or hypernyms sets, e.g., WordNet.

*e.g., synonym sets containing “good”:*

```
from nltk.corpus import wordnet as wn
poses = [ 'n': 'noun', 'v': 'verb', 's': 'adj (s)', 'a': 'adj', 'r': 'adv' ]
for synset in wn.synsets("good"):
    print("{}: {}".format(poses[synset.pos()],
        ", ".join([l.name() for l in synset.lemmas()])))
```

```
noun: good
noun: good, goodness
noun: good, goodness
noun: commodity, trade_good, good
adj: good
adj (sat): full, good
adj: good
adj (sat): estimable, good, honorable, respectable
adj (sat): beneficial, good
adj (sat): good
adj (sat): good, just, upright
...
adverb: well, good
adverb: thoroughly, soundly, good
```

*e.g., hypernyms of “panda”:*

```
from nltk.corpus import wordnet as wn
panda = wn.synset("panda.n.01")
hyper = lambda s: s.hypernyms()
list(panda.closure(hyper))
```

```
[Synset('procyonid.n.01'),
Synset('carnivore.n.01'),
Synset('placental.n.01'),
Synset('mammal.n.01'),
Synset('vertebrate.n.01'),
Synset('chordate.n.01'),
Synset('animal.n.01'),
Synset('organism.n.01'),
Synset('living_thing.n.01'),
Synset('whole.n.02'),
Synset('object.n.01'),
Synset('physical_entity.n.01'),
Synset('entity.n.01')]
```

# 1. How to Represent the Meaning of A Word?

**Issues** with WordNet-like solution.

- Missing nuance.
  - E.g., “proficient” is listed as a synonym for “good”.
  - This is only correct in certain contexts.
- Missing new meanings of words.
- Requires human labor to create and maintain.
  - Subjective.
- Cannot compute word similarity.
  - “good”, “great”, “proficient”, ...

## 2. Word Embeddings

- We still assign a **vector** to each word.
- The vector is **dense**, but not one-hot.

$$\text{banking} = \begin{pmatrix} 0.286 \\ 0.792 \\ -0.177 \\ -0.107 \\ 0.109 \\ -0.542 \\ 0.349 \\ 0.271 \end{pmatrix}$$

- The word vectors are also called word **embeddings** or (neural) word **representations**.

## 2. Word Embeddings

- **Word embeddings** – towards dense, semantically-meaningful representation of words.



Figure 1: Visualization of embeddings of words and phrases <sup>1</sup>.

<sup>1</sup><https://openclassrooms.com/en/courses/6532301-introduction-to-natural-language-processing/>

## 2. Word Embeddings

Why are they called “embeddings”?

“embed”: enroot, implant, ...



## 2. Word Embeddings

**The overall idea** to compute the embeddings of words.

- A word's meaning is decided by the words that frequently appear close-by.
  - “You shall know a word by the company it keeps” (J. R. Firth 1957: 11).
  - One of the most successful ideas of modern statistical NLP!
  - Representing words by their context.
- Given a word  $w$ , its context is the set of its nearby words (within a fixed-size window).





### 3. Word2Vec

Word2vec (Mikolov et al. 2013) is a computational framework for learning word embeddings.

- Input: A large corpus of texts.
- Output: Embeddings or word vectors. Every word in a fixed vocabulary is represented by a vector.
- Go through each position  $t$  in the text, which has a center word  $c$  and context (“outside”) words  $o$ .
- Calculate the probability of  $o$  given  $c$  (or vice versa).
- Adjust the vectors to maximize this probability.

### 3. Word2Vec




### 3. Word2Vec

For each position  $t = 1, \dots, T$ , predict context words within a window of fixed size  $m$  given the center word.

The likelihood of observing the given text:

$$\text{Likelihood} = L(\theta) = \prod_{t=1}^T \prod_{\substack{-m \leq j \leq m \\ j \neq 0}} P(w_{t+j} \mid w_t; \theta)$$

$\theta$  is all variables  
to be optimized



The objective function (cost function) is:

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} \mid w_t; \theta)$$

We want to minimize the objective function.

### 3. Word2Vec

We want to minimize the cost function:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

Question: How to calculate  $P(w_{t+j} | w_t; \theta)$ ?

Solution: We assign a vector  $\mathbf{v}_w$  to each word  $w$ .


Given a center word  $c$  and a context word  $o$ ,

$$P(o|c) = \frac{\exp(\mathbf{v}_o^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{v}_w^\top \mathbf{v}_c)}. \quad (2)$$

### 3. Word2Vec

Example of computing  $P(w_{t+j}|w_t; \theta)$

- $P(\text{turning}|\text{into}; \theta) = P(\mathbf{v}_{\text{turning}}|\mathbf{v}_{\text{into}})$



### 3. Word2Vec

A closer look at

$$P(o|c) = \frac{\exp(\mathbf{v}_o^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{v}_w^\top \mathbf{v}_c)} . \quad (3)$$

- This is an example of the **softmax function**  $\mathbb{R}^n \rightarrow (0, 1)^n$ .

$$\text{softmax}(x_i) = \frac{\exp(x_i)}{\sum_{j=1}^n \exp(x_j)} = p_i \quad (4)$$


- Maps arbitrary values  $x_i$  to a probability distribution  $p_i$  .
  - “max”: it amplifies the probability of largest  $x_i$ ;
  - “soft”: it still assigns some probability to smaller  $x_i$ .
  - A widely used technique in deep learning.

### 3. Word2Vec

Why is it a neural network?

### 3. Word2Vec

Why is it a neural network?





### 3. Word2Vec

Another look at


$$P(o|c) = \frac{\exp(\mathbf{v}_o^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{v}_w^\top \mathbf{v}_c)} . \quad (5)$$

- Computation cost?
- Solution: Negative sampling.

### 3. Word2Vec

- How to train our model?
  - Gradient descent.
  - Gradually adjust parameters, by walking down the gradient, to minimize the loss.
- What are the parameters  $\theta$ ?
  - Embedding vectors.
  - Each word owns two vectors.

$$\theta = \begin{bmatrix} v_{aardvark} \\ v_a \\ \vdots \\ v_{zebra} \\ u_{aardvark} \\ u_a \\ \vdots \\ u_{zebra} \end{bmatrix} \in \mathbb{R}^{2dV}$$



### 3. Word2Vec

#### Stochastic Gradient Descent


- $J(\theta)$  considers all windows (potentially billions) in the corpus.

$$J(\theta) = -\frac{1}{T} \log L(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

- Each parameter update is expensive.
- Solution: Stochastic gradient descent (SGD)
  - Repeatedly sample windows, and update part of the parameters for each mini-batch of windows.
  - Will it converge?

### 3. Word2Vec

$$\theta = \begin{bmatrix} v_{aardvark} \\ v_a \\ \vdots \\ v_{zebra} \\ u_{aardvark} \\ u_a \\ \vdots \\ u_{zebra} \end{bmatrix} \in \mathbb{R}^{2dV}$$



In SGD, for each mini-batch of data, the parameters update “somehow” moves towards the solution.

## 4. The “Real” Word2Vec

We want to minimize the cost function:

$$J(\theta) = -\frac{1}{T} \sum_{t=1}^T \sum_{\substack{-m \leq j \leq m \\ j \neq 0}} \log P(w_{t+j} | w_t; \theta)$$

When calculating  $P(w_{t+j} | w_t; \theta)$ , we assign two vectors to each word.

- $\mathbf{v}_w$ , when  $w$  is the center word,
- $\mathbf{u}_w$ , when  $w$  is the context word.

Given a center word  $c$  and a context word  $o$ ,

$$P(o|c) = \frac{\exp(\mathbf{u}_o^\top \mathbf{v}_c)}{\sum_{w \in \mathcal{V}} \exp(\mathbf{u}_w^\top \mathbf{v}_c)} . \quad (6)$$

Q: Which embedding to be used in downstream applications?

## 5. How to evaluate word embeddings?

- Intrinsic evaluation.
  - Evaluate on a specific/intermediate subtask.
  - Helps to understand the embeddings.
  - Limitations?
- Extrinsic evaluation.
  - Evaluate on a real task.
  - Text classification, QA, ...
  - Limitations?

## 5. Intrinsic evaluation

- Word Vector Analogies


a:b :: c:?

man:woman :: king:?



$$d = \arg \max_i \frac{(x_b - x_a + x_c)^T x_i}{\|x_b - x_a + x_c\|}$$

- Evaluate word vectors by how well their cosine distance after addition captures intuitive semantic and syntactic analogy questions
- Discarding the input words from the search!
- Problem: What if the information is there but not linear?



## 5. Intrinsic evaluation

