

# CSCI 4360/6360 Data Science II


## Recurrent Neural Networks

**Ninghao Liu**

Assistant Professor  
School of Computing  
University of Georgia

# In our previous lectures...


- The machine learning pipeline for text data.



- It is challenging to design feature vectors for text documents.
  - Documents can have arbitrary lengths.
  - How to encode semantic meanings?

# In our previous lectures...

- We know how to obtain **vector** representations for words.
- A word is represented by its **embedding**.
- A word embedding encodes the word's semantic meaning.
- A document contains multiple words.
- So how to represent the document with an embedding vector?



# Outline

- **Motivation**
  - High-level idea behind the new model architecture.
- **Major components of RNNs.**
- **RNN training.**
  - Supervised training
  - Unsupervised training

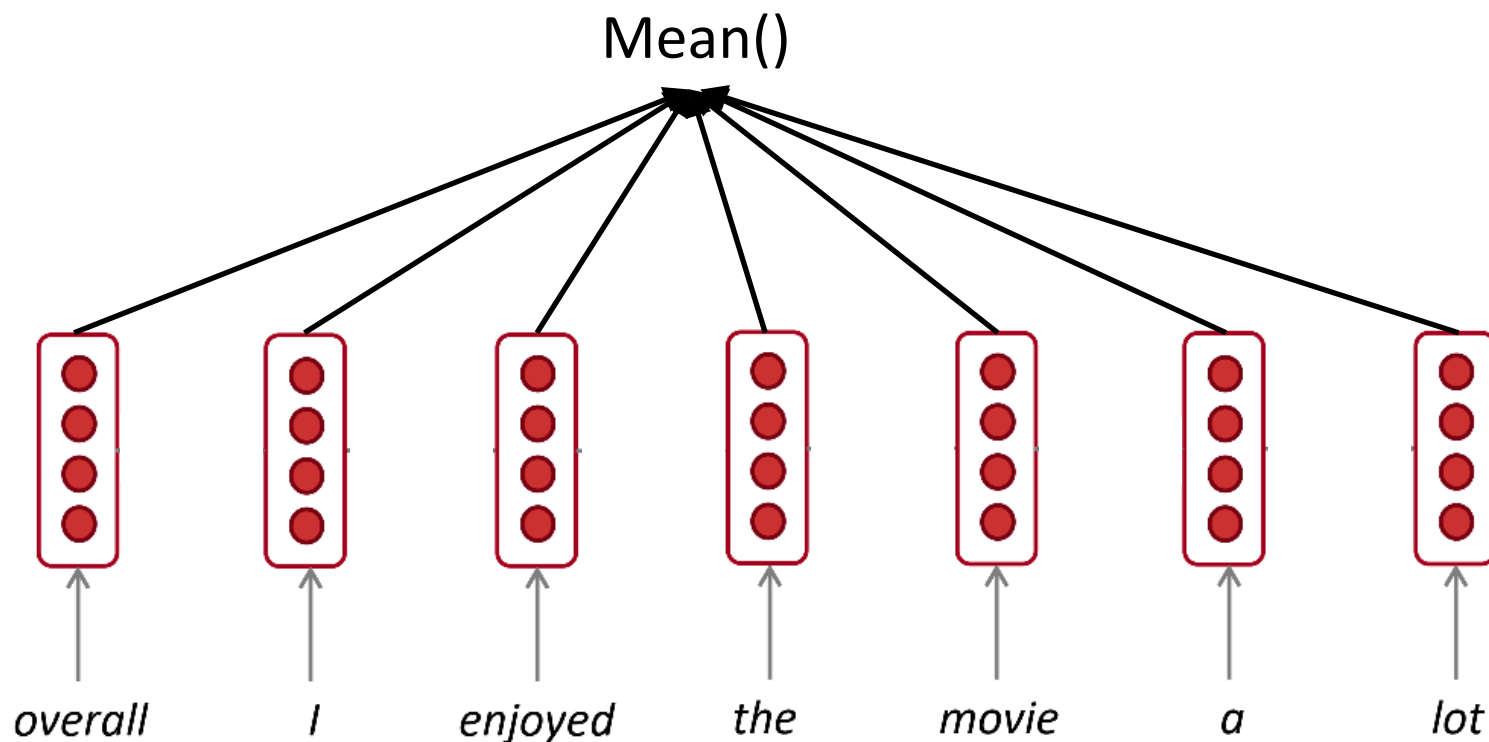
# Outline

- **Motivation**
  - High-level idea behind the new model architecture.
- **Major components of RNNs.**
- **RNN training.**
  - Supervised training
  - Unsupervised training

# Motivation

A native solution:

- Compute the **average** of word embeddings to represent the whole document.



# Motivation

## Challenge.

- Multiple words could be combined to express new ideas.
  - E.g., “I like the movie”, “I don’t like the movie”.
    - **Adding** the embeddings of “don’t” and “like” is not a good idea, because “don’t” contains no sentiment.
  - Many other examples
    - “check” + “in”  $\neq$  “check in”.
- To understand the logics, we treat sentences as **sequences**, where the order of words matters.

# Motivation

## Language Model

- The obvious question: how we should **model a sequence**?

What does it mean?

**Example 1:** Consider the following continuations of the phrase “It is raining -”.

1. “It is raining outside”
2. “It is raining banana tree”
3. “It is raining piouw;kcj pwepoiut”

**Example 2:** See blackboard.



# Motivation

## Language Model

- The obvious question: how we should **model a sequence**?
- Let  $(x_1, x_2, \dots, x_T)$  denote a sequence of words.
- By applying basic probability rules:

$$P(x_1, x_2, \dots, x_T) = \prod_{t=1}^T P(x_t \mid x_1, \dots, x_{t-1})$$

- For example,

$$\begin{aligned} &P(\text{deep, learning, is, fun}) \\ &= P(\text{deep})P(\text{learning} \mid \text{deep})P(\text{is} \mid \text{deep, learning})P(\text{fun} \mid \text{deep, learning, is}) \end{aligned}$$

# Motivation

## Language Model

- The probability is difficult to compute for **long** sequences:

$$P(x_1, x_2, \dots, x_T) = \prod_{t=1}^T P(x_t \mid x_1, \dots, x_{t-1})$$

- Thus, according to the Markov property, the impact of earlier words can be ignored.

$$\begin{aligned} P(x_1, x_2, x_3, x_4) &= P(x_1)P(x_2 \mid x_1)P(x_3 \mid x_2)P(x_4 \mid x_3), \\ P(x_1, x_2, x_3, x_4) &= P(x_1)P(x_2 \mid x_1)P(x_3 \mid x_1, x_2)P(x_4 \mid x_2, x_3). \end{aligned}$$

# Motivation

We need a novel model architecture that can

- Process sequential data.
- Store the information of **previous words** at each position.

**Recurrent neural networks (RNNs)** can handle this.

- But, instead of directly predicting  $P(x_t \mid x_1, \dots, x_{t-1})$ , RNN maintain a hidden state  $h_{t-1}$ , so that


$$P(x_t \mid x_{t-1}, \dots, x_1) \approx P(x_t \mid h_{t-1})$$

# Outline

- **Motivation**
  - High-level idea behind the new model architecture.
- **Major components of RNNs.**
- **RNN training.**
  - Supervised training
  - Unsupervised training

# RNN: Architecture (Simplified)

Core idea: Apply the same weights  $\mathbf{W}$  repeatedly



# RNN: Architecture

output distribution

$$\hat{y}^{(t)} = \text{softmax} \left( U h^{(t)} + b_2 \right) \in \mathbb{R}^{|V|}$$

hidden states

$$h^{(t)} = \sigma \left( W_h h^{(t-1)} + W_e e^{(t)} + b_1 \right)$$

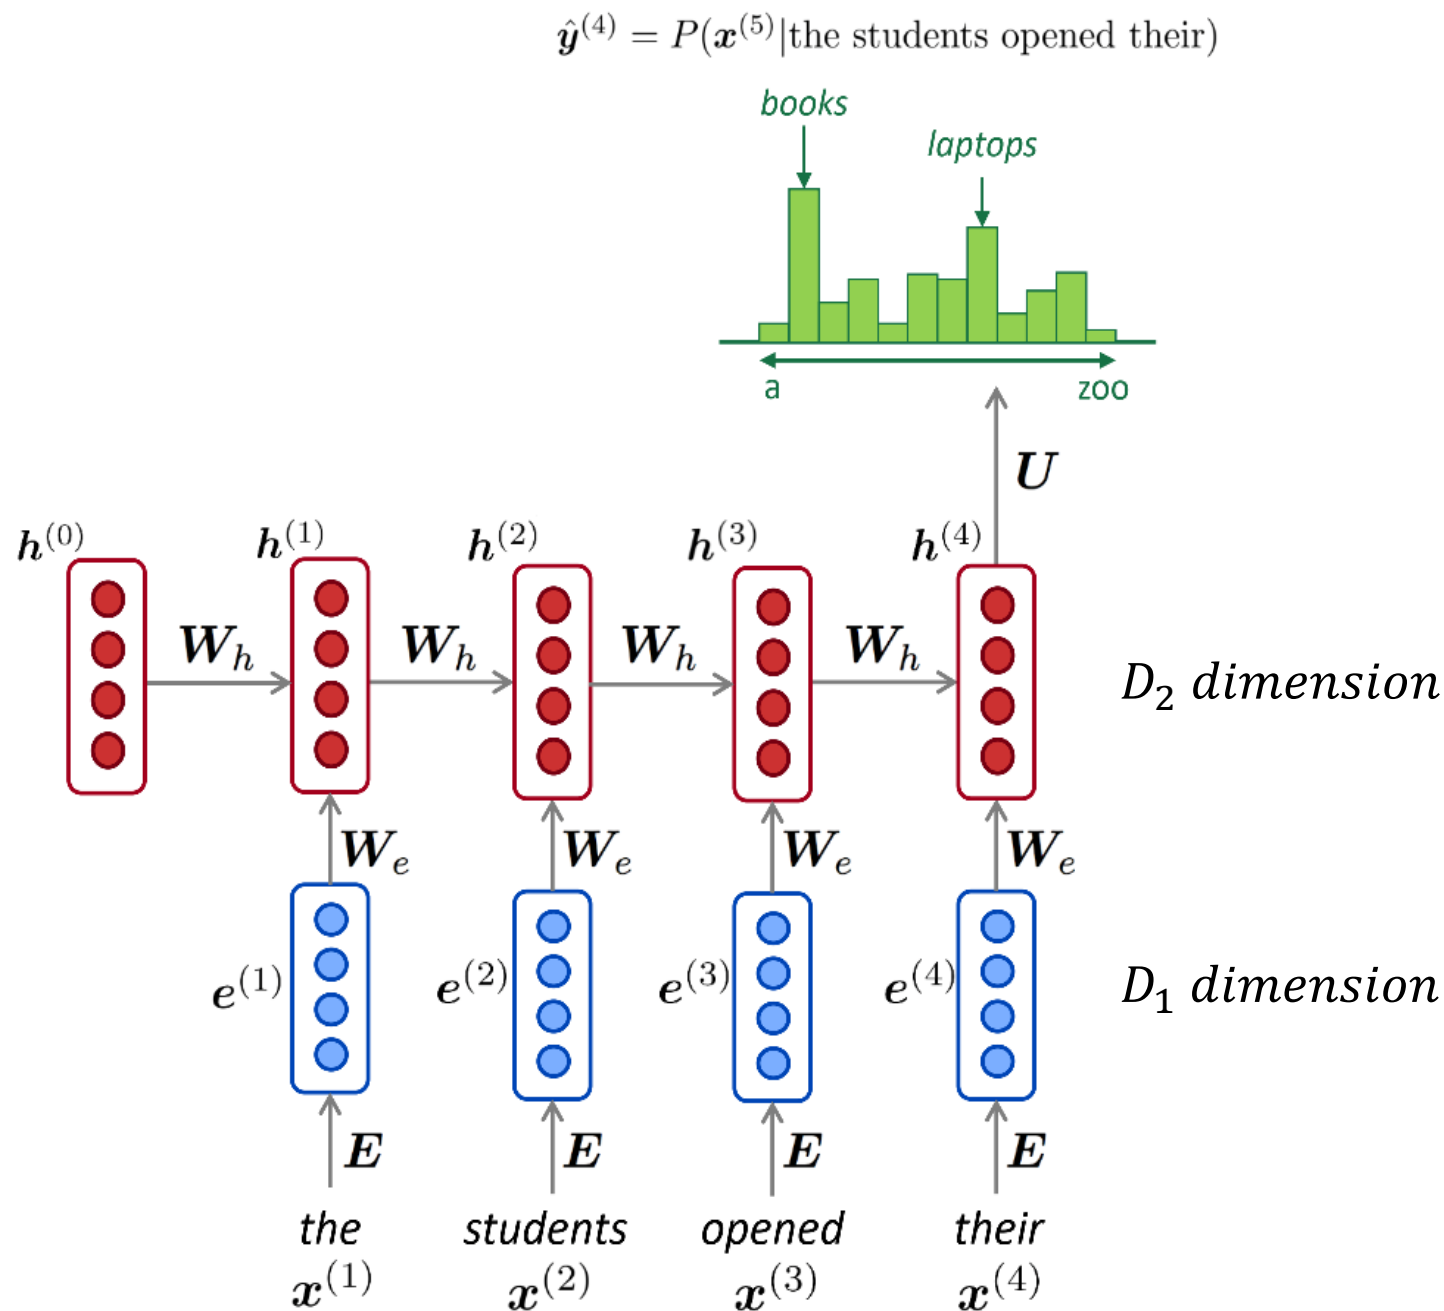
$h^{(0)}$  is the initial hidden state

word embeddings

$$e^{(t)} = E x^{(t)}$$

words / one-hot vectors


$$x^{(t)} \in \mathbb{R}^{|V|}$$



# RNN: Architecture

## RNN **Advantages:**

- Can process **any length** input
- Computation for step  $t$  can (in theory) use information from **many steps back**
- **Model size doesn't increase** for longer input
- Same weights applied on every timestep, so there is **symmetry** in how inputs are processed.




# RNN: Architecture

## RNN **Disadvantages:**

- Recurrent computation is **slow**
- In practice, difficult to access information from **many steps back**

More on these later in the course

RNN is designed to process sequential data, but it cannot handle very-long sequences.





# Outline

- **Motivation**
  - High-level idea behind the new model architecture.
- **Major components of RNNs.**
- **RNN training.**
  - Supervised training
  - Unsupervised training

# RNN: Training

- **Supervised learning**

- The data is labeled

- $D = \{(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_N, y_N)\}$

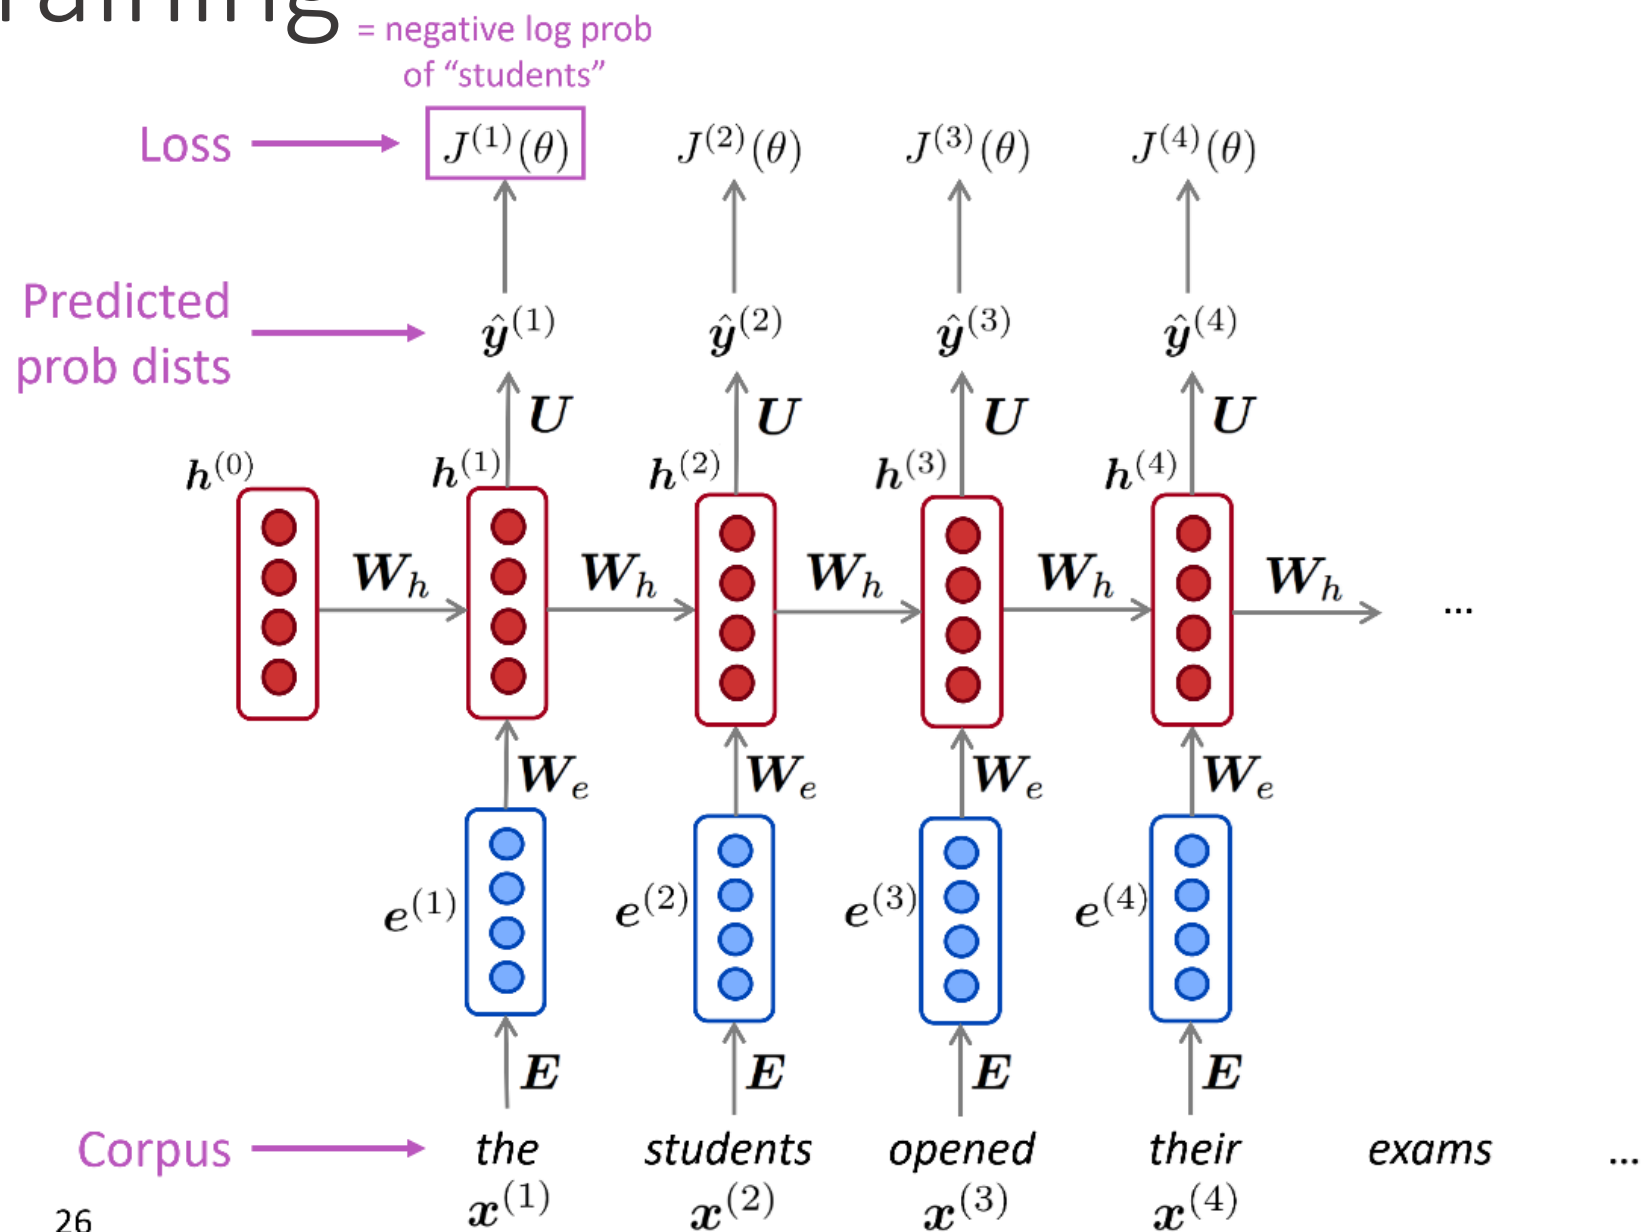
- Procedure:

- Let RNNs read the whole text, and **produce an output**
- Apply a **loss function** to measure the degree of match between model output and the label.
- Similar to other deep models, such as **MLPs**.

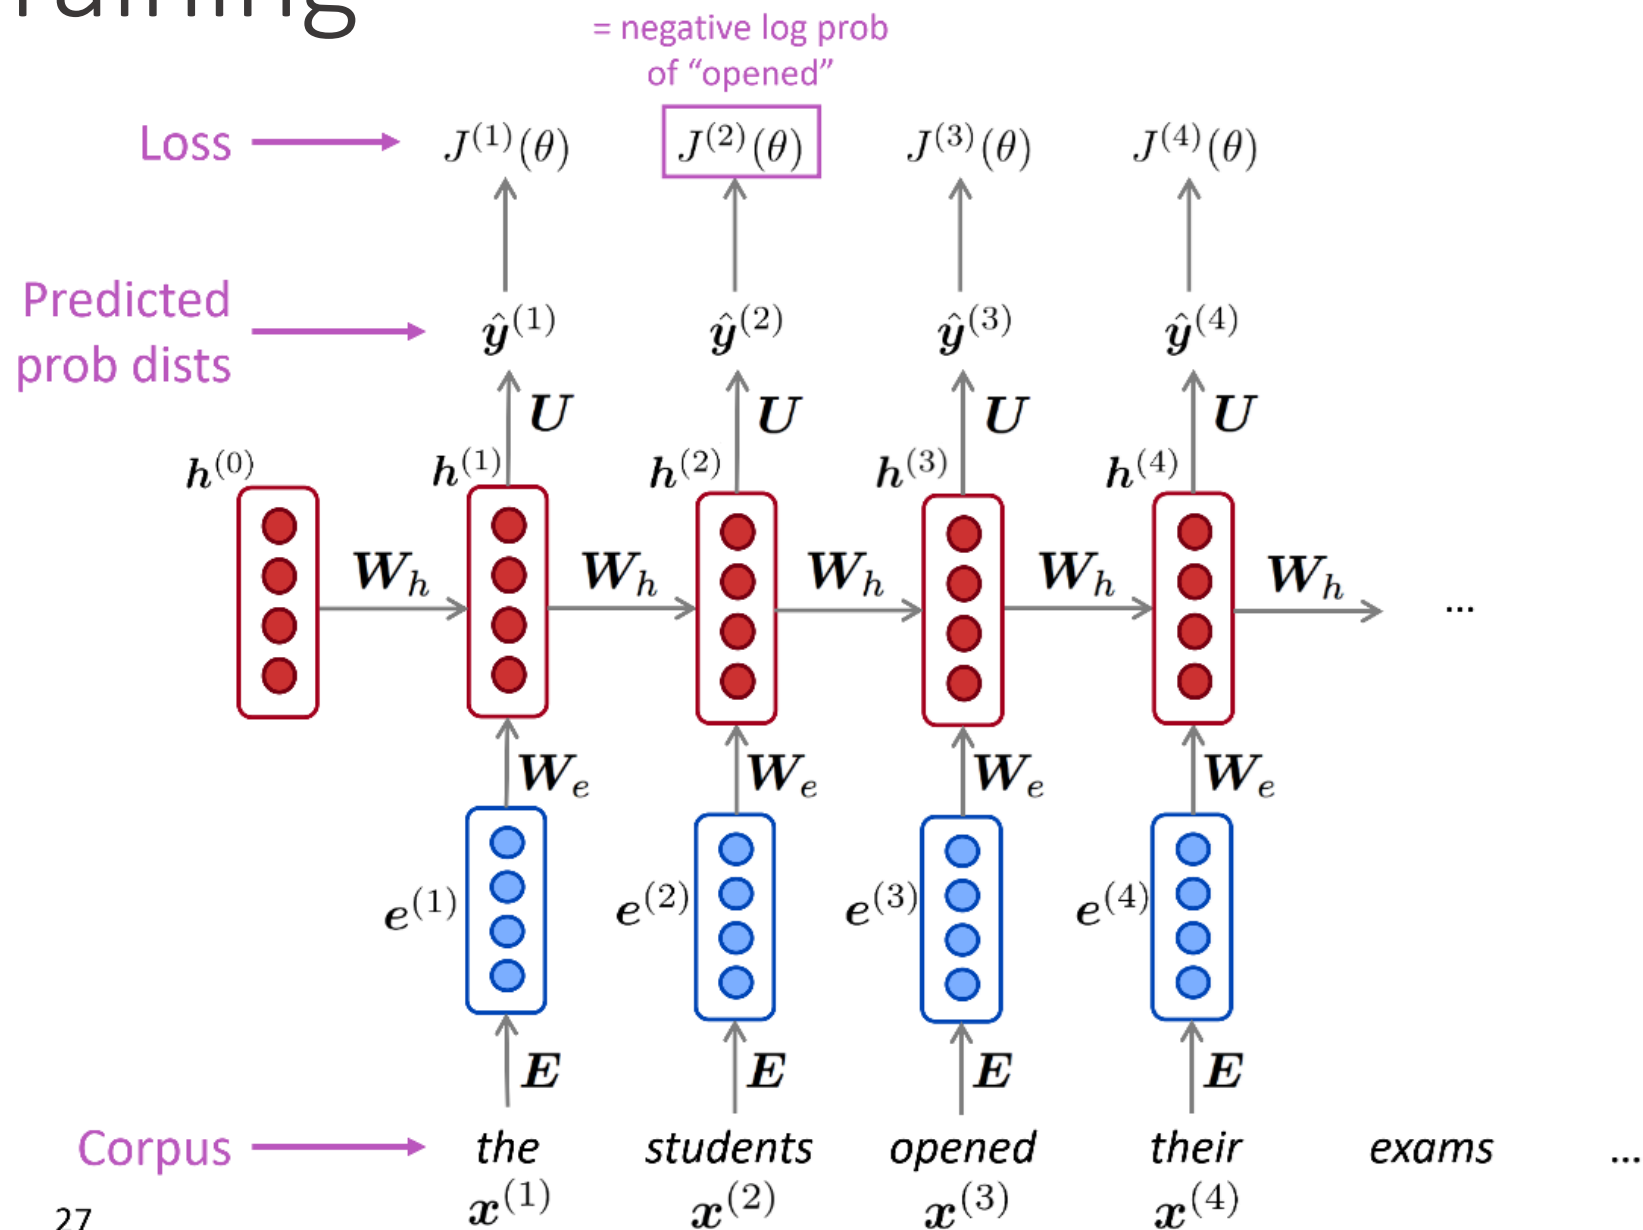
# RNN: Training

- RNNs (and many other NLP models) can be trained with **unsupervised learning**.
  - $D = \{x_1, x_2, \dots, x_N\}$

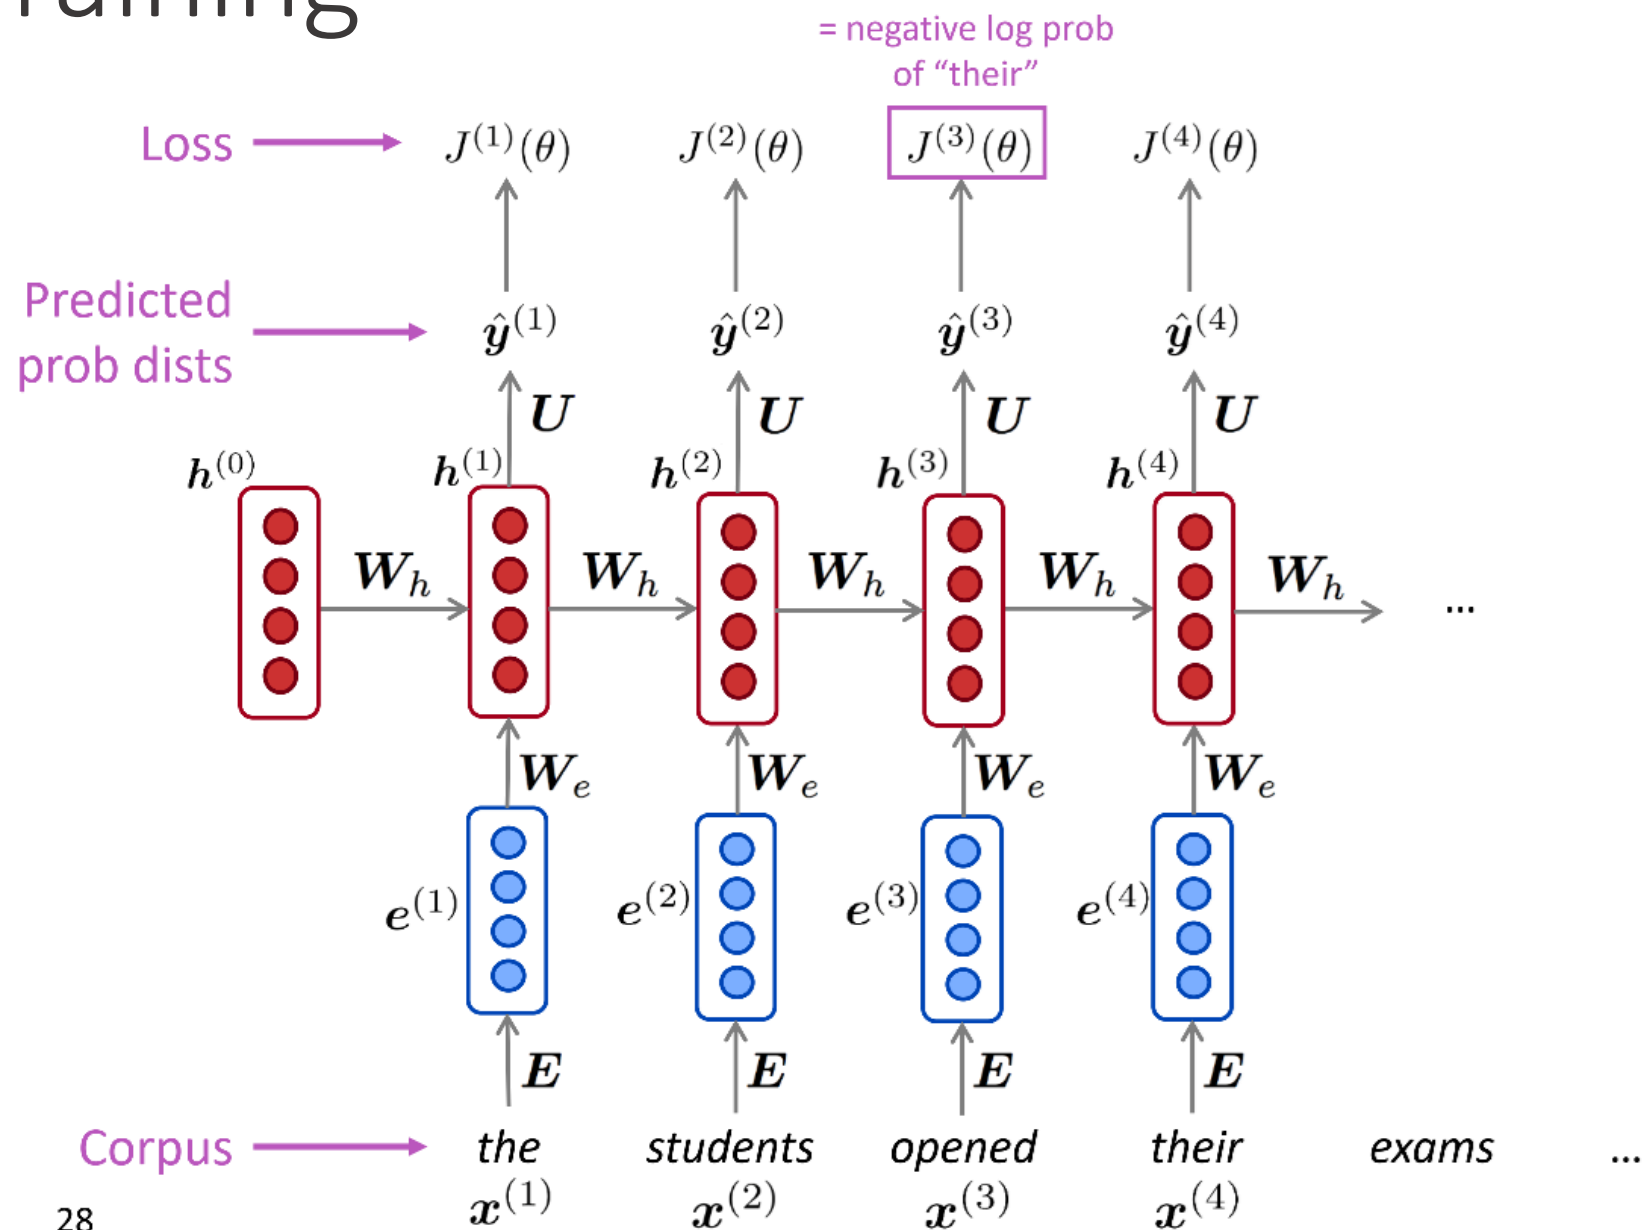
# RNN: Training



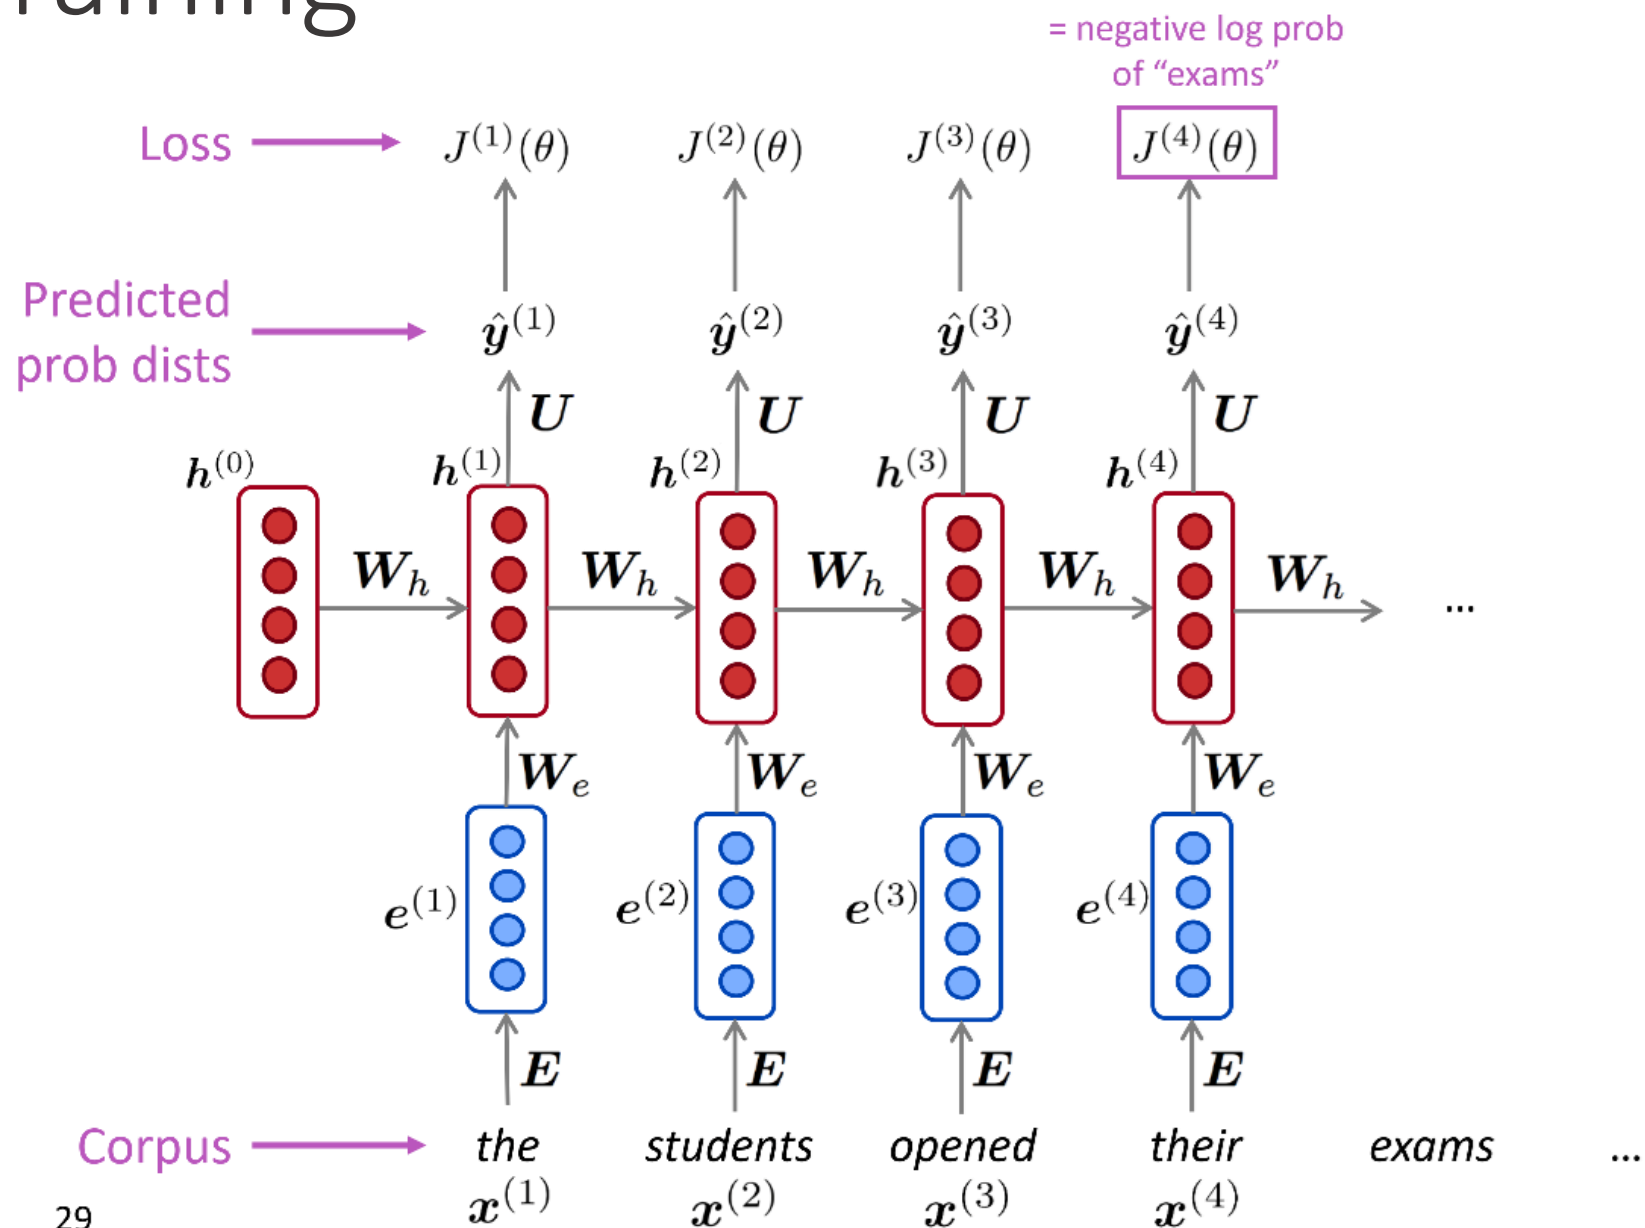
# RNN: Training




# RNN: Training



# RNN: Training



# RNN: Training





# RNN: Training


- However: Computing loss and gradients across **entire corpus**  $x^{(1)}, \dots, x^{(T)}$  is **too expensive!**

$$J(\theta) = \frac{1}{T} \sum_{t=1}^T J^{(t)}(\theta)$$

- In practice, consider  $x^{(1)}, \dots, x^{(T)}$  as a **sentence** (or a **document**)
- Recall: **Stochastic Gradient Descent** allows us to compute loss and gradients for small chunk of data, and update.
- Compute loss  $J(\theta)$  for a sentence (actually a batch of sentences), compute gradients and update weights. Repeat.

# RNN: Training - Example

- **Goal:** Train an RNN model to do time-series prediction. (code available on eLC.)



# RNN: Training - Example

- **Step 1:** Import *numpy* and *pytorch* modules

```
import torch
from torch import nn
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline
```

# RNN: Training

- **Step 2: Define RNN.**

Define model components.



Define data flow.



```
class RNN(nn.Module):
```

```
    def __init__(self, input_size, output_size, hidden_dim, n_layers):
        super(RNN, self).__init__()

        self.hidden_dim=hidden_dim

        # define an RNN with specified parameters
        # "batch_first": the first dim of input and output is the batch_size
        self.rnn = nn.RNN(input_size, hidden_dim,
                           n_layers, batch_first=True)

        # last, fully-connected layer
        self.fc = nn.Linear(hidden_dim, output_size)
```

```
    def forward(self, x, hidden):
        # x (batch_size, seq_length, input_size)
        # hidden (n_layers, batch_size, hidden_dim)
        # r_out (batch_size, time_step, hidden_size)
        batch_size = x.size(0)


        # get RNN outputs
        r_out, hidden = self.rnn(x, hidden)
        # shape output to be (batch_size*seq_length, hidden_dim)
        r_out = r_out.view(-1, self.hidden_dim)

        # get final output
        output = self.fc(r_out)

        return output, hidden
```

# RNN: Training

- **Step 2.1: RNN components**



```
class RNN(nn.Module):
    def __init__(self, input_size, output_size, hidden_dim, n_layers):
        super(RNN, self).__init__()

        self.hidden_dim=hidden_dim


        # define an RNN with specified parameters
        # "batch_first": the first dim of input and output is the batch_size
        self.rnn = nn.RNN(input_size, hidden_dim,
                           n_layers, batch_first=True)

        # last, fully-connected layer
        self.fc = nn.Linear(hidden_dim, output_size)
```

# RNN: Training

- **Step 2.2: RNN flow**

```
def forward(self, x, hidden):  
    # x (batch_size, seq_length, input_size)  
    # hidden (n_layers, batch_size, hidden_dim)  
    # r_out (batch_size, time_step, hidden_size)  
    batch_size = x.size(0)  
  
    # get RNN outputs  
    r_out, hidden = self.rnn(x, hidden)  
    # shape output to be (batch_size*seq_length, hidden_dim)  
    r_out = r_out.view(-1, self.hidden_dim)  
  
    # get final output  
    output = self.fc(r_out)  
  
    return output, hidden
```



# RNN: Training

- **Step 3: Training Loss and Optimizer**

```
# MSE loss and Adam optimizer with a learning rate of 0.01
criterion = nn.MSELoss()
optimizer = torch.optim.Adam(rnn.parameters(), lr=0.01)
```

# RNN: Training

- **Step 4: Training Process**

Prepare data for this batch

Feedforward flow

Gradient descent

```
# train the RNN
def train(rnn, n_steps, print_every):

    # initialize the hidden state
    hidden = None

    for batch_i, step in enumerate(range(n_steps)):

        # defining the training data
        time_steps = np.linspace(step * np.pi, (step+1)*np.pi, seq_length + 1)
        data = np.sin(time_steps)
        data.resize((seq_length + 1, 1)) # input_size=1

        x = data[:-1]
        y = data[1:]

        # convert data into Tensors
        x_tensor = torch.Tensor(x).unsqueeze(0) # unsqueeze gives a 1, batch_size dimension
        y_tensor = torch.Tensor(y)

        # outputs from the rnn
        prediction, hidden = rnn(x_tensor, hidden)

        ## Representing Memory ##
        # make a new variable for hidden and detach the hidden state from its history
        # this way, we don't backpropagate through the entire history
        hidden = hidden.data

        # calculate the loss
        loss = criterion(prediction, y_tensor)
        # zero gradients
        optimizer.zero_grad()
        # perform backprop and update weights
        loss.backward()
        optimizer.step()

        # display loss and predictions
        if batch_i % print_every == 0:
            print('Loss: ', loss.item())
            plt.plot(time_steps[1:], x, 'r.') # input
            plt.plot(time_steps[1:], prediction.data.numpy().flatten(), 'b.') # predictions
            plt.show()

    return rnn
```



# RNN: Training

- **Step 4.1: Get training data for this batch**


```
# defining the training data
time_steps = np.linspace(step * np.pi, (step+1)*np.pi, seq_length + 1)
data = np.sin(time_steps)
data.resize((seq_length + 1, 1)) # input_size=1

x = data[:-1]
y = data[1:]

# convert data into Tensors
x_tensor = torch.Tensor(x).unsqueeze(0) # unsqueeze gives a 1, batch_size dimension
y_tensor = torch.Tensor(y)
```

# RNN: Training

- **Step 4.1: Get training data for this batch (if visualized)**



# RNN: Training

- **Step 4.2: Feedforward flow**

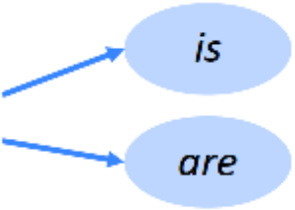


Call **forward()**

```
# outputs from the rnn  
prediction, hidden = rnn(x_tensor, hidden)
```

```
def forward(self, x, hidden):  
    # x (batch_size, seq_length, input_size)  
    # hidden (n_layers, batch_size, hidden_dim)  
    # r_out (batch_size, time_step, hidden_size)  
    batch_size = x.size(0)  
  
    # get RNN outputs  
    r_out, hidden = self.rnn(x, hidden)  
    # shape output to be (batch_size*seq_length, hidden_dim)  
    r_out = r_out.view(-1, self.hidden_dim)  
  
    # get final output  
    output = self.fc(r_out)  
  
    return output, hidden
```

# RNN: Training - Vanishing Gradient

## Effect of vanishing gradient on RNN-LM

- **LM task:** *The writer of the books \_\_\_\_* 
- **Correct answer:** *The writer of the books is planning a sequel*
- **Syntactic recency:** *The writer of the books is* (correct) 
- **Sequential recency:** *The writer of the books are* (incorrect) 
- Due to vanishing gradient, RNN-LMs are better at learning from **sequential recency** than **syntactic recency**, so they make this type of error more often than we'd like [Linzen et al 2016]

# RNN: Training - Vanishing Gradient

- **Problems** with RNNs!
  - Vanishing gradients

motivates

- **Fancy RNN** variants!
  - LSTM
  - GRU
  - multi-layer
  - bidirectional

People are not satisfied!

Transformer