

# CSCI 4360/6360 Data Science II

## Transformers


**Ninghao Liu**

Assistant Professor  
School of Computing  
University of Georgia

# Where we left off

## Self-Attention:

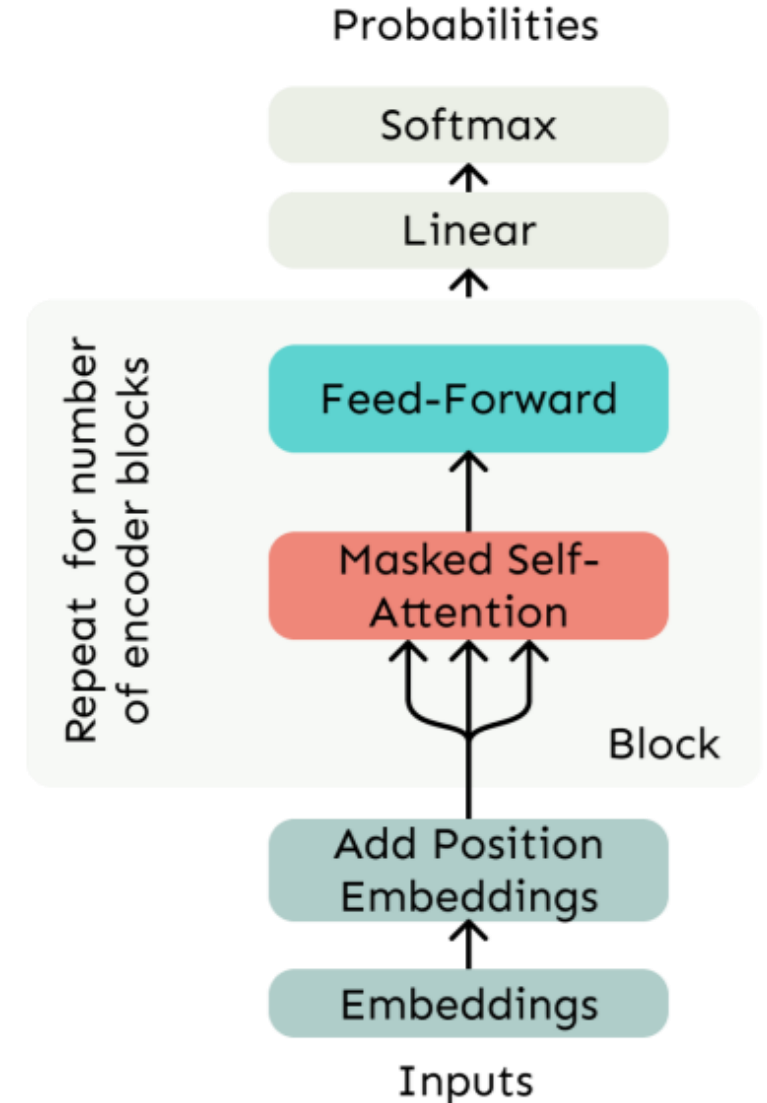
- Words attend to themselves.
- A word plays the role of query, key, and value simultaneously.



All words attend to all words in previous layer; most arrows here are omitted

# Where we left off

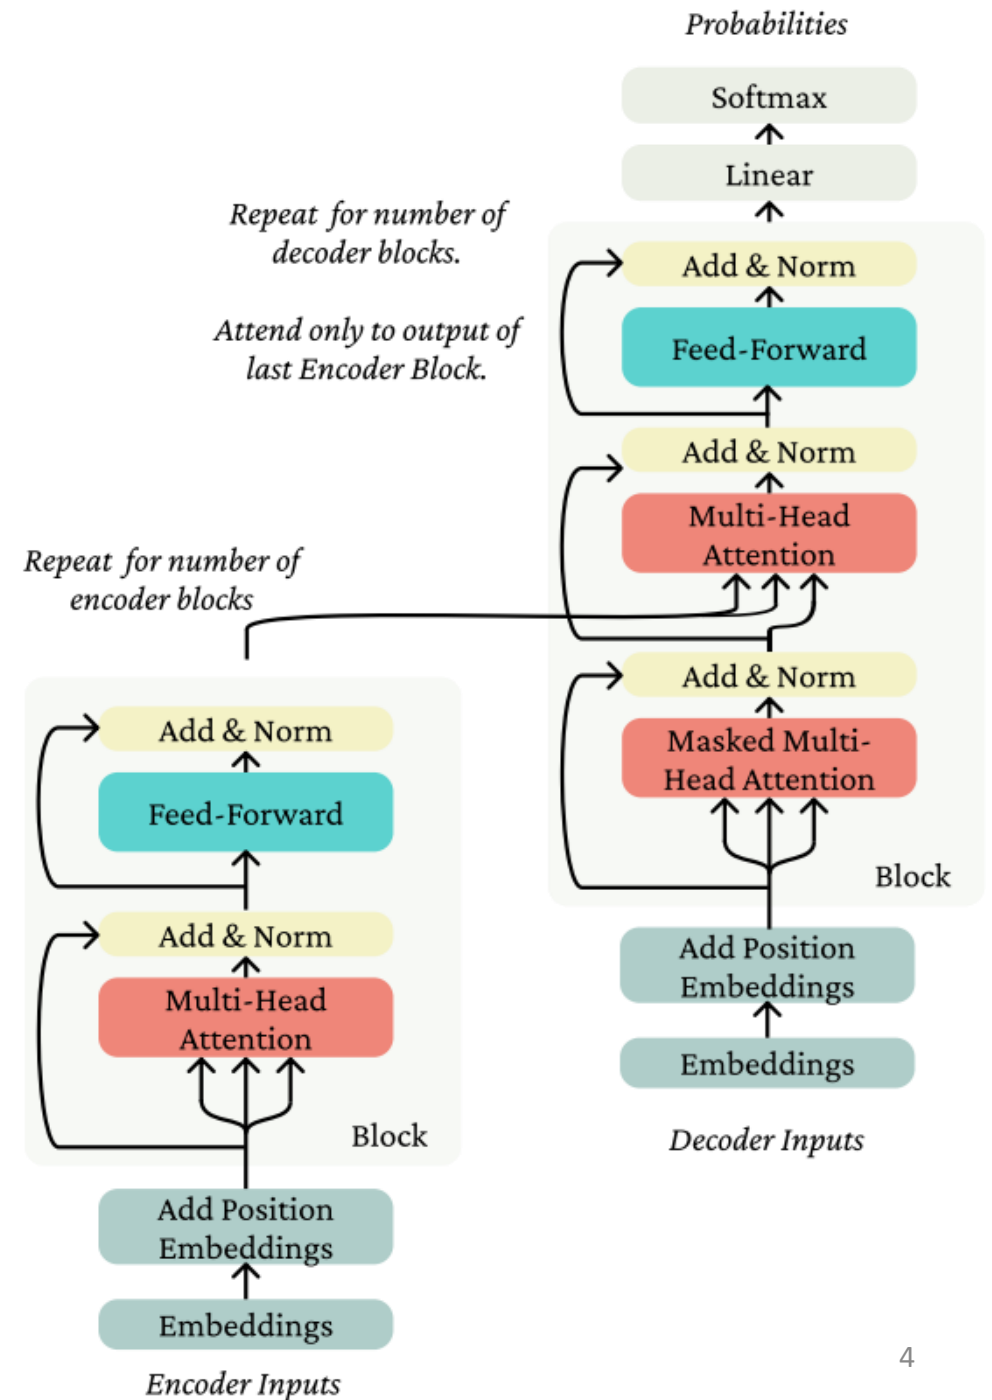
- **Self-attention:**
  - the basis of the method.
- **Position representations:**
  - Specify the sequence order, since self-attention is an unordered function of its inputs.
- **Nonlinearities:**
  - At the output of the self-attention block
  - Frequently implemented as a simple feed-forward network.
- **Masking:**
  - In order to parallelize operations while not looking at the future.
  - Keeps information about the future from “leaking” to the past.



# Transformer - Overview

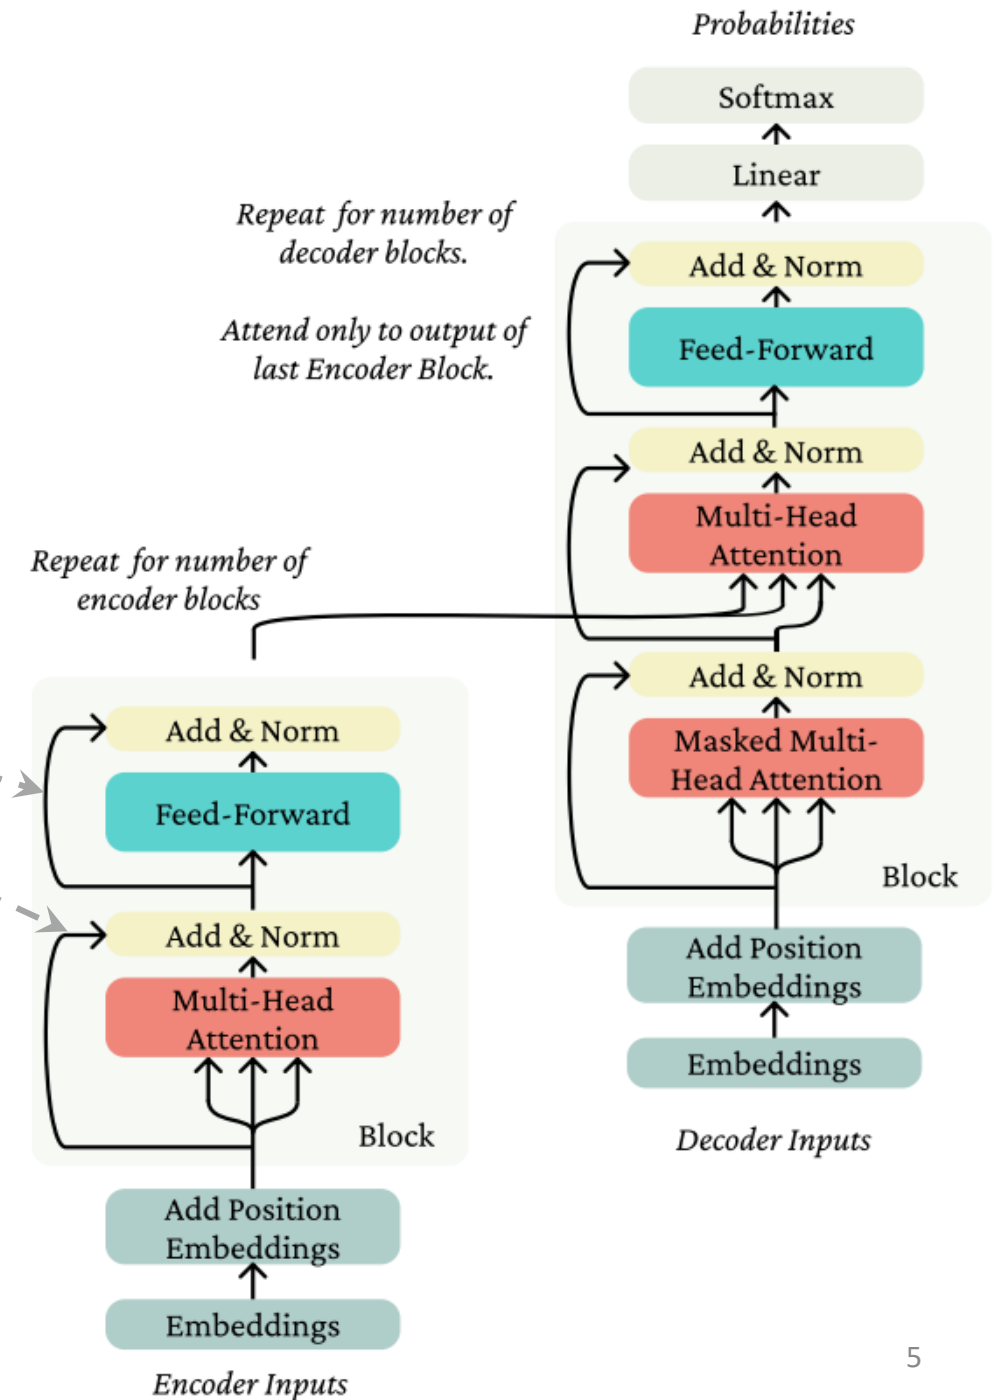
## Key points

- Originally, an Encoder-Decoder architecture
- Self-Attention Block
  - Self-attention mechanism
  - Position embeddings
  - Nonlinearity in FF layers
  - Masking (optional)



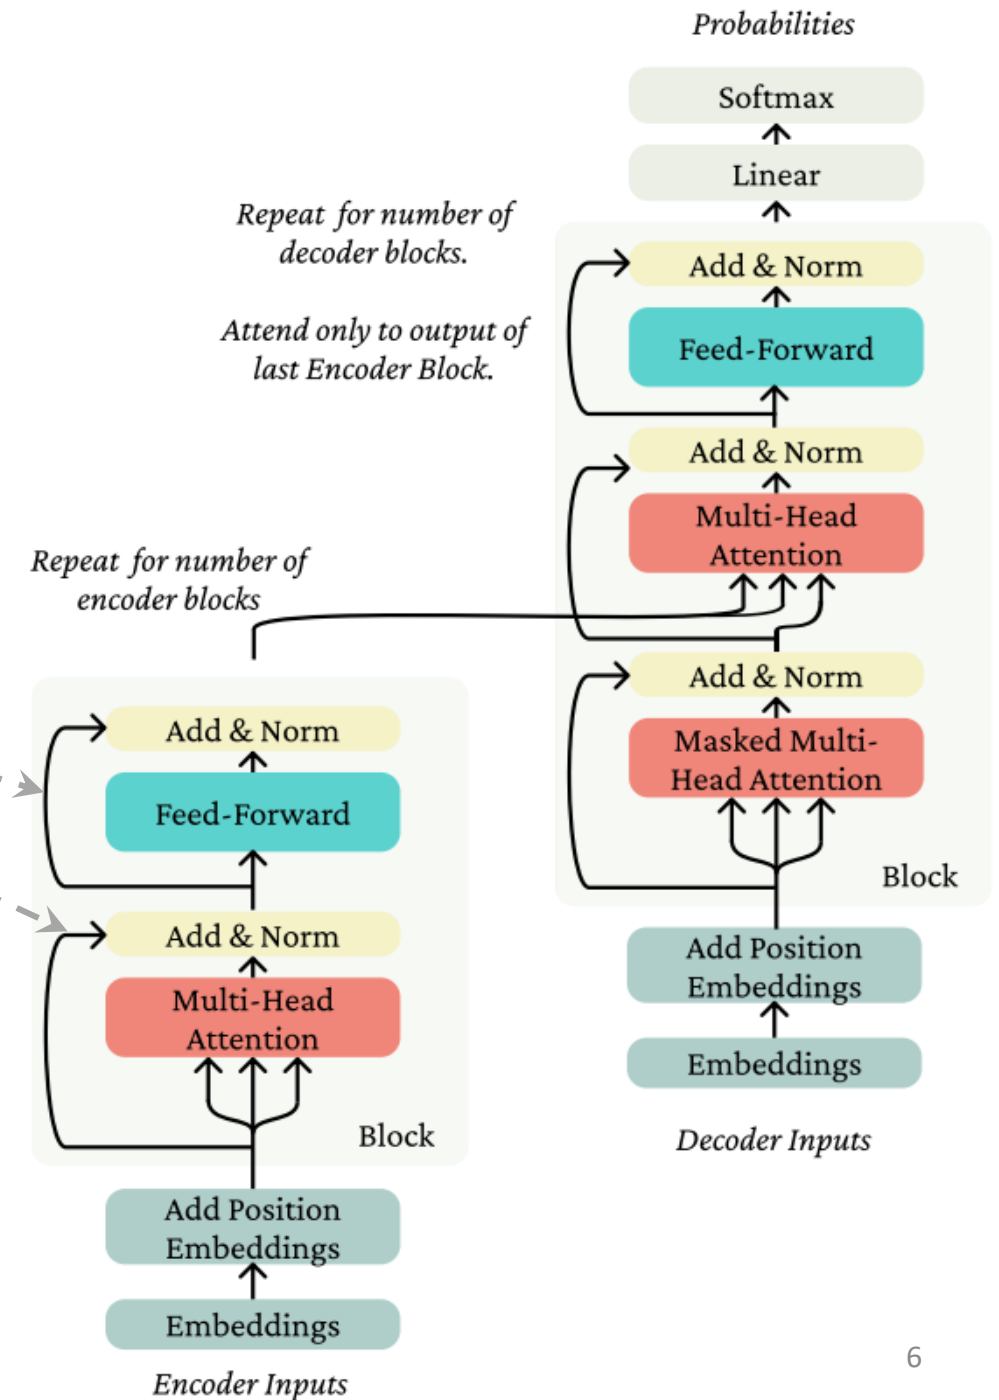
# Transformer - Overview

- What is “Multi-Head Attention”?
- What is this?
- What is “Norm”



# Transformer - Overview

- **What is “Multi-Head Attention”?**
- What is this?
- What is “Norm”




# Multi-Head Self-Attention

- Multi-Head Self-Attention
  - “Self-Attention” + “Multi-Head”


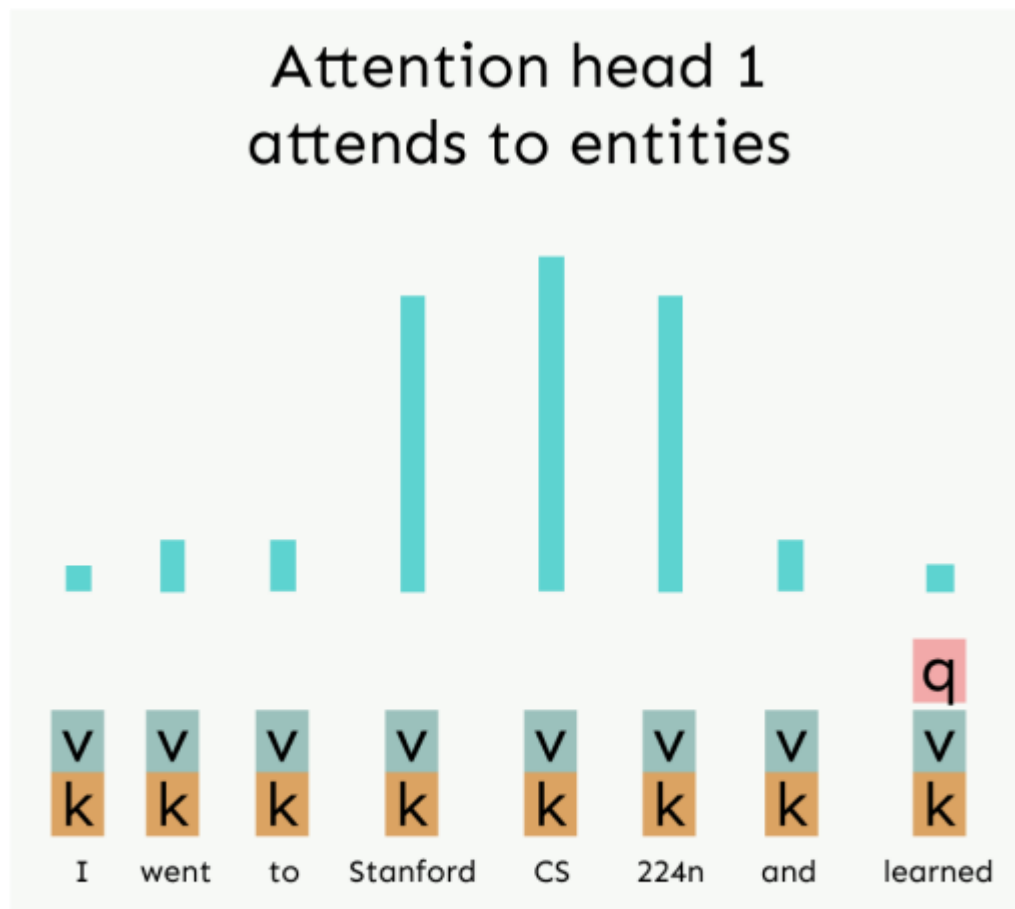
# Multi-Head Self-Attention

## Recall the Self-Attention Hypothetical Example





# Multi-Head Self-Attention



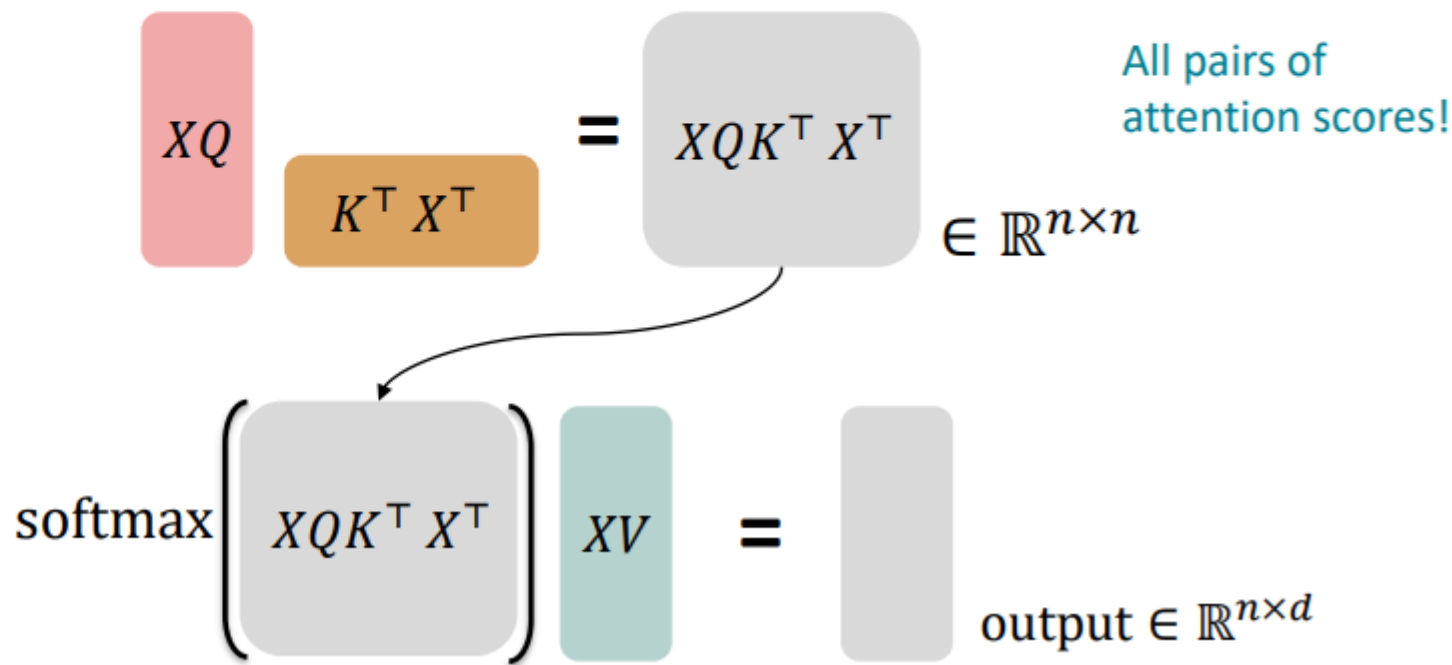
I    went    to    Stanford

CS    224n    and    learned

# Single-Head Self-Attention

- Let's look at how key-query-value attention is computed, in matrices.
  - Let  $X = [x_1; \dots; x_n] \in \mathbb{R}^{n \times d}$  be the concatenation of input vectors.
  - First, note that  $XK \in \mathbb{R}^{n \times d}$ ,  $XQ \in \mathbb{R}^{n \times d}$ ,  $XV \in \mathbb{R}^{n \times d}$ .
  - The output is defined as  $\text{output} = \text{softmax}(XQ(XK)^T)XV \in \mathbb{R}^{n \times d}$ .

First, take the query-key dot products in one matrix multiplication:  $XQ(XK)^T$



Next, softmax, and compute the weighted average with another matrix multiplication.

# Multi-Head Self-Attention

- What if we want to look in multiple places in the sentence at once?
  - For word  $i$ , self-attention “looks” where  $x_i^\top Q^\top K x_j$  is high, but maybe we want to focus on different  $j$  for different reasons?
- We’ll define **multiple attention “heads”** through multiple Q,K,V matrices
- Let,  $Q_\ell, K_\ell, V_\ell \in \mathbb{R}^{d \times \frac{d}{h}}$ , where  $h$  is the number of attention heads, and  $\ell$  ranges from 1 to  $h$ .
- Each attention head performs attention independently:
  - $\text{output}_\ell = \text{softmax}(X Q_\ell K_\ell^\top X^\top) * X V_\ell$ , where  $\text{output}_\ell \in \mathbb{R}^{d/h}$
- Then the outputs of all the heads are combined!
  - $\text{output} = [\text{output}_1; \dots; \text{output}_h] Y$ , where  $Y \in \mathbb{R}^{d \times d}$
- Each head gets to “look” at different things, and construct value vectors differently.

# Multi-Head Self-Attention

- Even though we compute  $h$  many attention heads, it's not really more costly.
  - We compute  $XQ \in \mathbb{R}^{n \times d}$ , and then reshape to  $\mathbb{R}^{n \times h \times d/h}$ . (Likewise for  $XK, XV$ .)
  - Then we transpose to  $\mathbb{R}^{h \times n \times d/h}$ ; now the head axis is like a batch axis.
  - Almost everything else is identical, and the **matrices are the same sizes**.

First, take the query-key dot products in one matrix multiplication:  $XQ(XK)^\top$

The diagram illustrates the first step of Multi-Head Self-Attention. On the left, a red vertical bar represents the query matrix  $XQ$ . In the middle, an orange horizontal bar represents the product  $K^\top X^\top$ . An equals sign follows, leading to a gray rounded rectangle representing the resulting matrix  $XQK^\top X^\top$ . To the right of this rectangle is the text  $\in \mathbb{R}^{3 \times n \times n}$ . Above the gray rectangle, the text "3 sets of all pairs of attention scores!" is written in blue. A curved line connects the gray rectangle to the "softmax" operation in the diagram below.

Next, softmax, and compute the weighted average with another matrix multiplication.

The diagram illustrates the second step of Multi-Head Self-Attention. It starts with the text "softmax" followed by a large left square bracket. Inside the bracket is a gray rounded rectangle labeled  $XQK^\top X^\top$ . To the right of the bracket is a teal vertical bar labeled  $XV$ . An equals sign follows, leading to a gray vertical bar. To its right is a small gray rounded rectangle labeled  $P$  with the word "mix" underneath it. Another equals sign follows, leading to a final gray vertical bar. To the right of this bar is the text "output  $\in \mathbb{R}^{n \times d}$ ".

# Scaled Dot Product

- **“Scaled Dot Product”** attention aids in training.
- When dimensionality  $d$  becomes large, dot products between vectors tend to become large.
  - Because of this, inputs to the softmax function can be large, making the gradients small.

- Instead of the self-attention function we’ve seen:

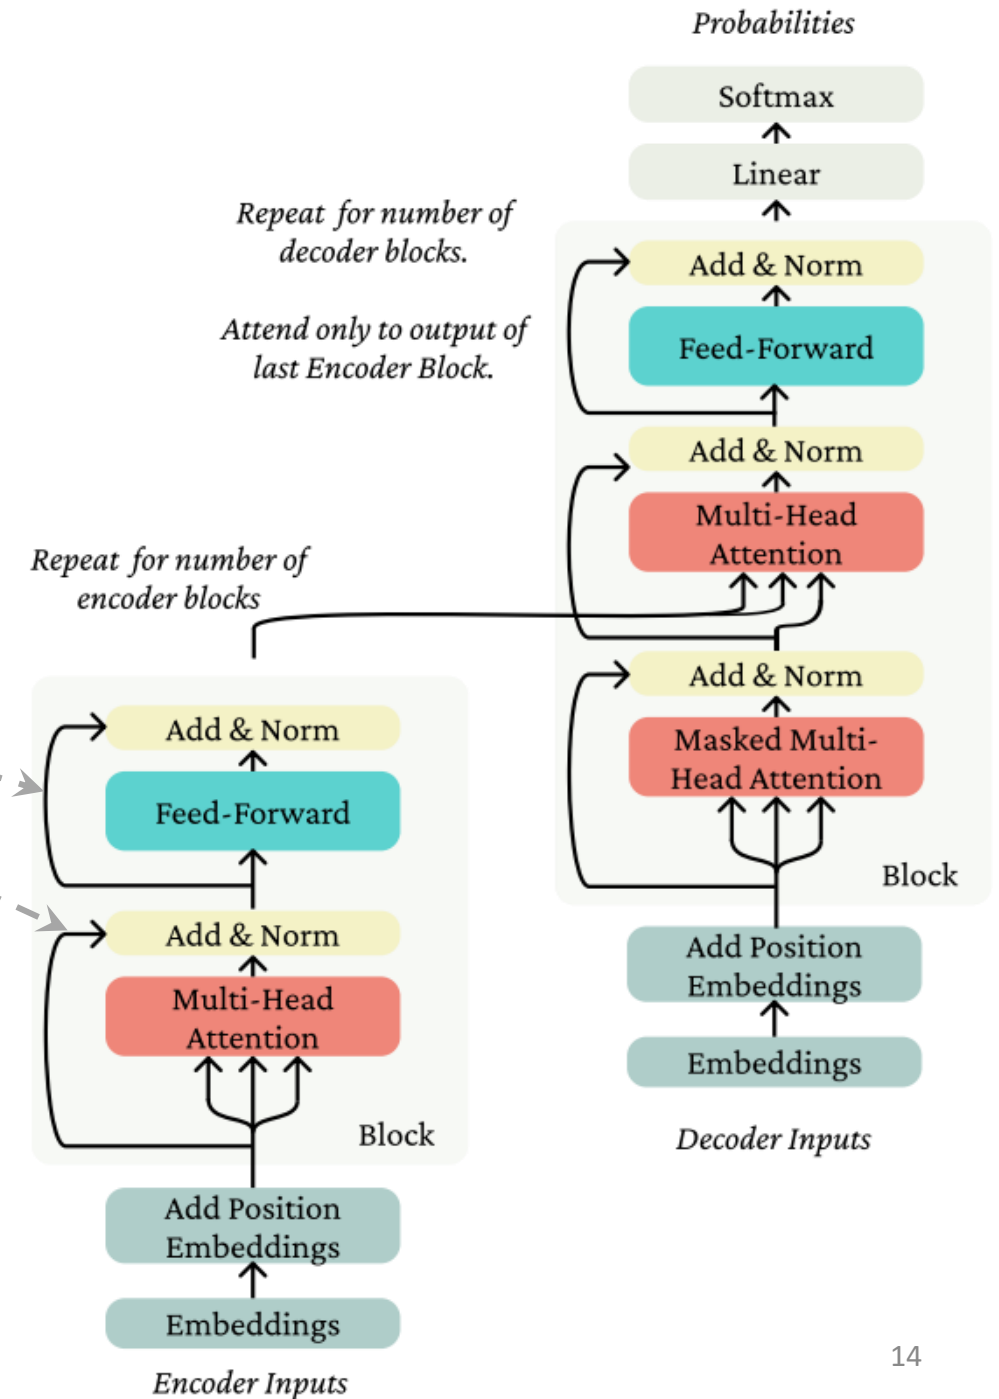
$$\text{output}_\ell = \text{softmax}(XQ_\ell K_\ell^\top X^\top) * XV_\ell$$

- We divide the attention scores by  $\sqrt{d/h}$ , to stop the scores from becoming large just as a function of  $d/h$  (The dimensionality divided by the number of heads.)

$$\text{output}_\ell = \text{softmax}\left(\frac{XQ_\ell K_\ell^\top X^\top}{\sqrt{d/h}}\right) * XV_\ell$$

# Transformer - Overview

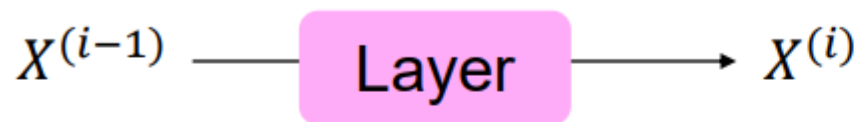
- What is “Multi-Head Attention”?
- **What is this?**
- What is “Norm”






# Residual Connection


- **Residual connections** are a trick to help models train better.
  - Instead of  $X^{(i)} = \text{Layer}(X^{(i-1)})$  (where  $i$  represents the layer)




- We let  $X^{(i)} = X^{(i-1)} + \text{Layer}(X^{(i-1)})$  (so we only have to learn “the residual” from the previous layer)



- Gradient is **great** through the residual connection; it's 1!
- Bias towards the identity function!



[no residuals]

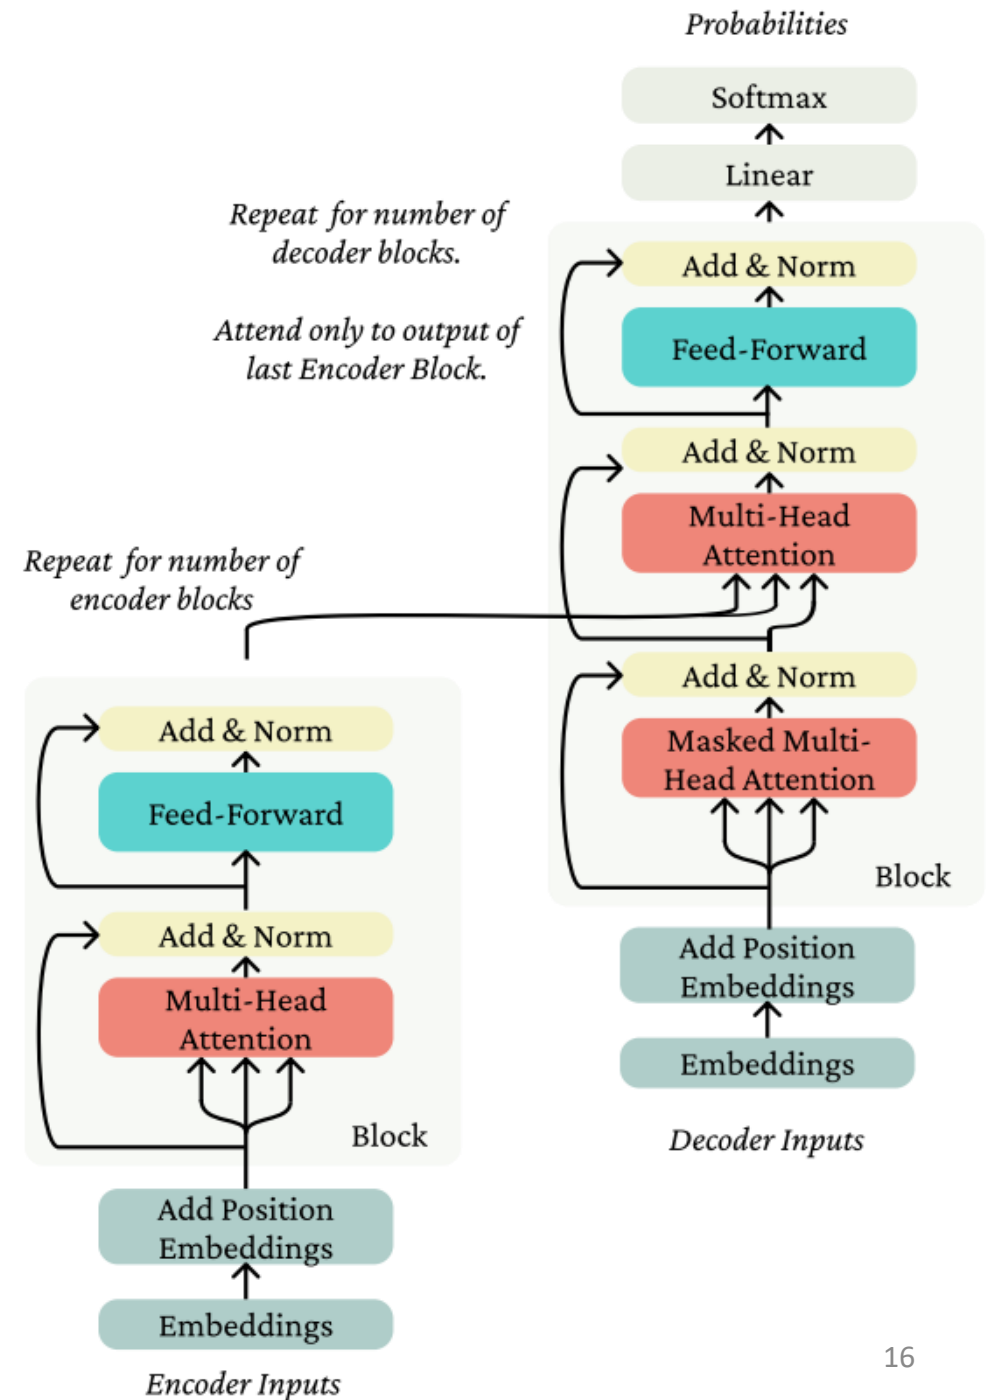


[residuals]

[Loss landscape visualization,  
[Li et al., 2018](#), on a ResNet]

# Transformer - Overview

- What is “Multi-Head Attention”?
- What is this?
- **What is “Norm”**





# Layer Normalization

- **Layer normalization** is a trick to help models train faster.
- Idea: cut down on uninformative variation in hidden vector values by normalizing to unit mean and standard deviation **within each layer**.
  - LayerNorm's success may be due to its normalizing gradients [[Xu et al., 2019](#)]
- Let  $x \in \mathbb{R}^d$  be an individual (word) vector in the model.
- Let  $\mu = \sum_{j=1}^d x_j$ ; this is the mean;  $\mu \in \mathbb{R}$ .
- Let  $\sigma = \sqrt{\frac{1}{d} \sum_{j=1}^d (x_j - \mu)^2}$ ; this is the standard deviation;  $\sigma \in \mathbb{R}$ .
- Let  $\gamma \in \mathbb{R}^d$  and  $\beta \in \mathbb{R}^d$  be learned “gain” and “bias” parameters. (Can omit!)
- Then layer normalization computes:

$$\text{output} = \frac{x - \mu}{\sqrt{\sigma} + \epsilon} * \gamma + \beta$$

Normalize by scalar mean and variance

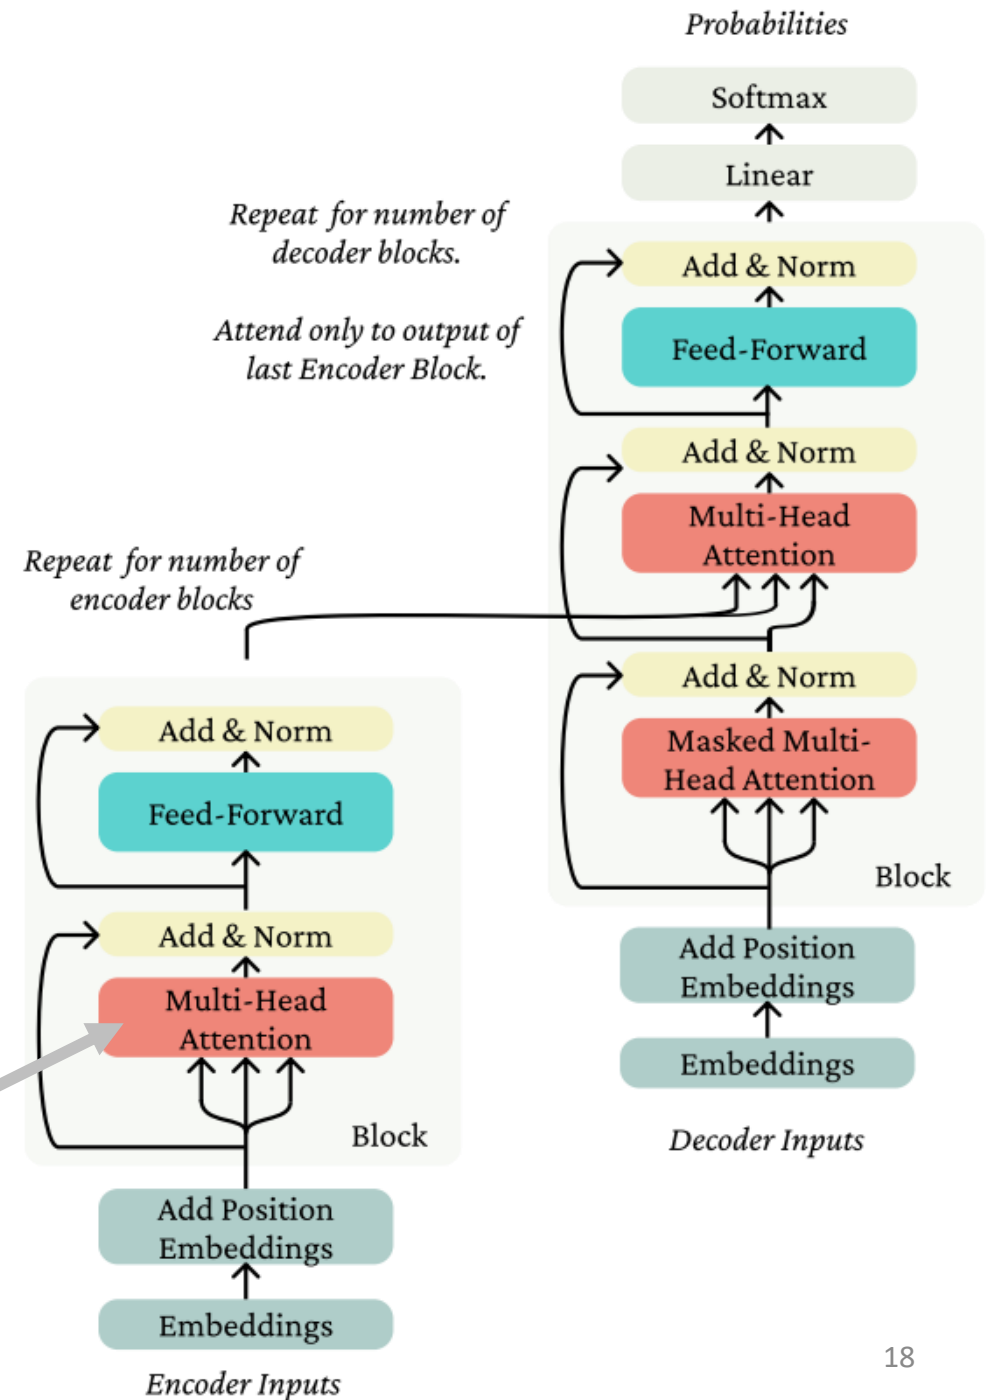
Modulate by learned elementwise gain and bias

# Transformer – Wrap Up

## Key points

- Encoder - Decoder architecture
- But! the encoder and decoder can be used **separately**!

No Masking!

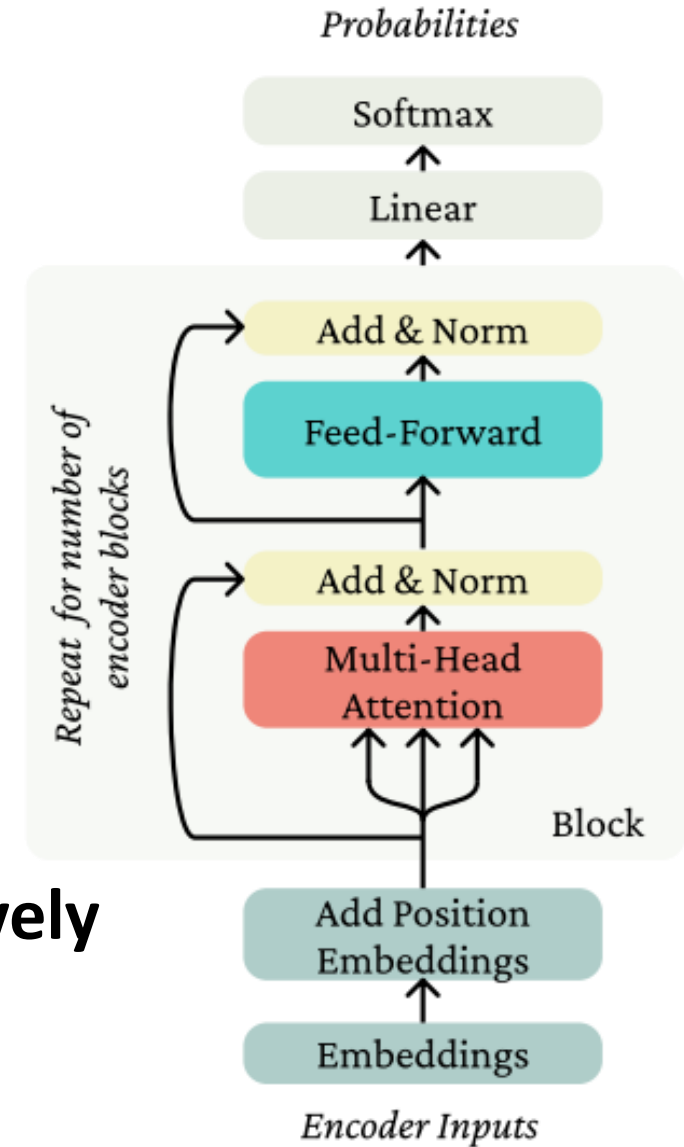


# Transformer – Wrap Up

- Each Block consists of:
  - Self-attention
  - Add & Norm
  - Feed-Forward
  - Add & Norm

A Transformer Encoder is great when:

- You **aren't** trying to generate text **autoregressively** (there's no masking in the encoder so each position index can see the whole sequence,);
- Famous examples: BERT, RoBERTa.



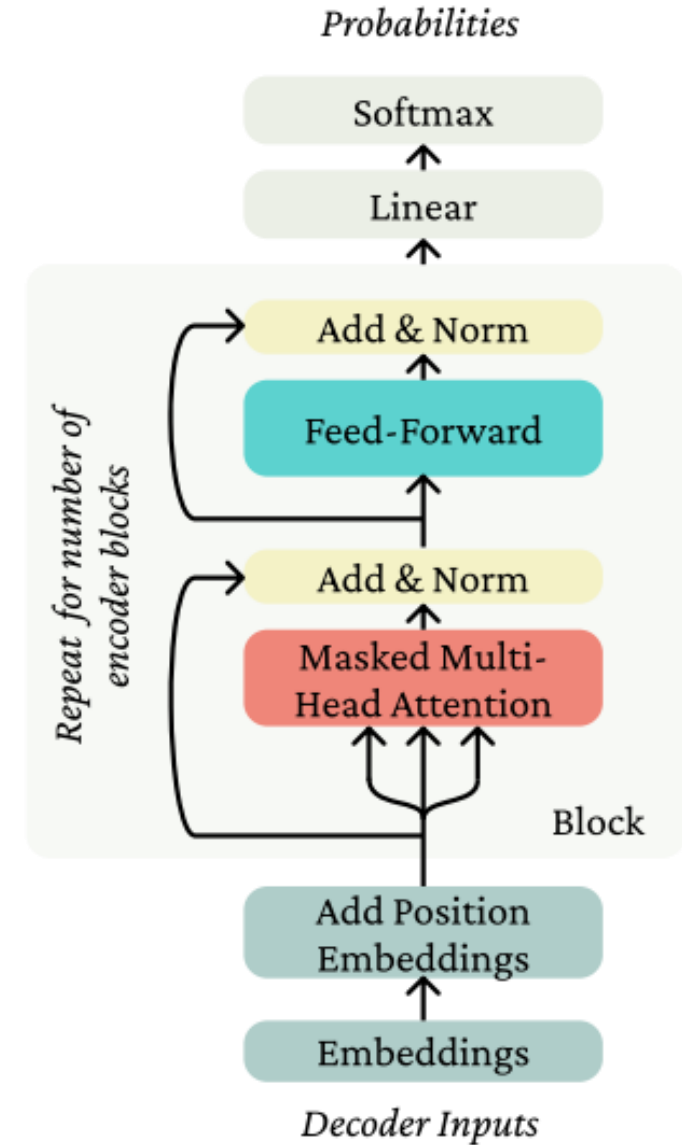
*Transformer Encoder*

# Transformer – Wrap Up

Decoder is very similar to Encoder.

- To build an **autoregressive language model**, one uses a Transformer Decoder.
- Using future masking at each application of self-attention.

Famous example: GPT-2, GPT-3, BLOOM.



*Transformer Decoder*

# Transformer – Wrap Up

When to use the whole Encoder - Decoder architecture?

- When we'd like bidirectional context on something (e.g., **article summarization**, **translation**)
- Such an architecture has been found to provide better performance than decoder-only models

