

Outlier Detection

Ninghao Liu

University of Georgia

April 10, 2024

Overview

- 1 Background and Problem Definition
- 2 Outlier Detection (OD) Strategies
 - Supervised, Unsupervised, Semi-supervised
- 3 Outlier Detection Methods
 - {Statistical, Proximity, Clustering, Classification} based
 - k-means for Outlier Detection
 - Isolation Forest
- 4 Dealing with High-Dimensional Data in OD
 - Autoencoders for OD

1. A Real-World Application

Imagine you are a transaction auditor in a credit card company:

- Protect customers from credit card fraud.
- Pay special attention to card usages that are different from typical cases.
- Not-that-typical cases: the purchase occurs *far* from the owner's resident city, the purchase amount is much *bigger* than usual, very *frequent* purchases in a short time, ...
- Detect such transactions as soon as they occur and contact the card owner for verification.

1. Problem Definition

An **outlier** is a data object that deviates significantly from the rest of the objects, as if it were generated by a different mechanism.





Figure 1: The objects in region R are outliers.

1. Problem Definition

Outlier detection (or **anomaly detection**) is the process of finding data objects with behaviors that are very different from expectation.

Please note that:

- Outliers are different from *noisy* data.
- It may be important to justify *why* the outliers are detected.
- Outlier detection is also related to *novelty detection* (but they are not the same).



2. Outlier Detection Strategies

Outlier detector training strategies:

- Supervised Learning
- Unsupervised Learning
- Semi-supervised Learning
- ...

2. Outlier Detection Strategies

2.1 Supervised Learning for OD

Given a dataset $\{(\mathbf{x}_i, y_i)\}$, $1 \leq i \leq N$, train a classifier $f : \mathcal{X} \rightarrow \{0, 1\}$. For a new instance \mathbf{x} , it is classified as an outlier if $f(\mathbf{x}) = 1$, or a normal instance if $f(\mathbf{x}) = 0$.

2. Outlier Detection Strategies

2.1 Supervised Learning for OD

Given a dataset $\{(\mathbf{x}_i, y_i)\}$, $1 \leq i \leq N$, train a classifier $f : \mathcal{X} \rightarrow \{0, 1\}$. For a new instance \mathbf{x} , it is classified as an outlier if $f(\mathbf{x}) = 1$, or a normal instance if $f(\mathbf{x}) = 0$.

Challenges:

- Data imbalance: Methods are needed for handling data imbalance.
- Evaluation gap: Classification maximizes prediction accuracy, which may not be of interest in applications.
- Labeling difficulty: The sampled data for labeling may not sufficiently represent the outlier distribution.

2. Outlier Detection Strategies

2.2 Unsupervised Learning for OD

Given a dataset $\{\mathbf{x}_i\}$, $1 \leq i \leq N$, train a detector $f : \mathcal{X} \rightarrow \{0, 1\}$. For a new instance \mathbf{x} , it is regarded as an outlier if $f(\mathbf{x}) = 1$, or a normal instance if $f(\mathbf{x}) = 0$.

2. Outlier Detection Strategies

2.2 Unsupervised Learning for OD

Given a dataset $\{\mathbf{x}_i\}$, $1 \leq i \leq N$, train a detector $f : \mathcal{X} \rightarrow \{0, 1\}$. For a new instance \mathbf{x} , it is regarded as an outlier if $f(\mathbf{x}) = 1$, or a normal instance if $f(\mathbf{x}) = 0$.

- Implicit assumption: The normal objects are “clustered”.
- This assumption may not be true in some cases.
 - Outliers may also be “clustered”.
 - Normal objects are diversely distributed.
- Computational challenge: We have to process a large population of non-target data entries (i.e., the normal objects) before one can touch the real meat (i.e., the outliers).

2. Outlier Detection Strategies




Figure 2: The black objects form outliers collectively.

2. Outlier Detection Strategies

2.3 Semi-supervised Learning for OD

Given a dataset $\{\mathbf{x}_i\}$ and $\{y_j\}$, $1 \leq i \leq N, j \in \mathcal{J}, |\mathcal{J}| \ll N$, train a detector $f : \mathcal{X} \rightarrow \{0, 1\}$.

- Most of the instances are unlabeled, so effective unsupervised methods serve as the important basis.
- Both labeled normal/outlier instances are possible.

3. Outlier Detection Methods

- Statistical Methods
- Clustering-Based Methods
- Proximity-Based Methods
- ...

3.1 Statistical Methods

The general idea: Learn a probability density function $p(\mathbf{x})$, denoting the probability that \mathbf{x} is generated by the normal mechanism.

- Parametric methods: Assume p follows certain priori distribution (e.g., gaussian) with parameters θ to be estimated.
- Non-parametric methods: Do not assume a priori distribution.

3.1 Statistical Methods

Steps of parametric methods:

- Find a certain distribution to work with.
- Estimate the parameters θ of the distribution.
- Calculate $p(\mathbf{x}; \theta)$, and decide if the probability is low enough.

Example 1:

Univariate outlier detection using maximum likelihood. Suppose a city's average temperature values in July in the last 10 years are, in value-ascending order, 24.0°C, 28.9°C, 28.9°C, 29.0°C, 29.1°C, 29.1°C, 29.2°C, 29.2°C, 29.3°C, and 29.4°C.

$$p(\mathbf{x}) = \frac{1}{\sigma\sqrt{2\pi}} \cdot \exp\left(-\frac{1}{2} \frac{(x-\mu)^2}{\sigma^2}\right)$$


3.1 Statistical Methods

Steps of parametric methods:

- Find a certain distribution to work with.
- Estimate the parameters θ of the distribution.
- Calculate $p(\mathbf{x}; \theta)$, and decide if the probability is low enough.

Example 2:

$$p(\mathbf{x}) = (2\pi)^{-\frac{k}{2}} \cdot \det(\Sigma)^{-\frac{1}{2}} \cdot \exp\left(-\frac{1}{2}(\mathbf{x} - \boldsymbol{\mu})^\top \Sigma^{-1}(\mathbf{x} - \boldsymbol{\mu})\right)$$




3.1 Statistical Methods

Commonly used non-parametric techniques:

- Histogram
- Kernel based density estimation

Example:

Outlier detection using a histogram. *AllElectronics* records the purchase amount for every customer transaction. Figure 12.5 uses a histogram (refer to Chapters 2 and 3) to graph these amounts as percentages, given all transactions.



Q: Should a transaction amounted \$6500 be regarded as an outlier?

3.1 Statistical Methods

Steps of histogram methods:

- Build the histogram from data.
- Given an instance x , check it against the histogram, and assign it with an outlier score of $\frac{1}{hist(x)}$.

Challenges


- How to choose the bin size?

3.1 Statistical Methods

Another Method: Treat an observed object as an indicator of high probability density in the surrounding region.

3.1 Statistical Methods

Another Method¹: Treat an observed object as an indicator of high probability density in the surrounding region.



¹<https://blogs.sas.com/content/iml/2016/07/27/visualize-kernel-density-estimate.html>

3.1 Statistical Methods

Kernel density estimation

Given samples $\{\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N\}$ and a test instance \mathbf{x} ,

$$p(\mathbf{x}) = \frac{1}{Nh} \sum_{i=1}^N K\left(\frac{\mathbf{x} - \mathbf{x}_i}{h}\right) \quad (1)$$

where $K()$ is the kernel function, h is the bandwidth parameter.

A kernel $K()$ is a non-negative real-valued function that satisfies:

- $\int_{-\infty}^{+\infty} K(u) du = 1$,
- $K(u) = K(-u)$.

A frequently used kernel is the RBF kernel: $K(\mathbf{x}) = \exp(-\gamma \cdot \|\mathbf{x}\|_2^2)$.

3.1 Statistical Methods

Properties of statistical methods:

- Statistically justifiable
- Handling **high-dimensional data** is challenging
- Computational cost: For simple parametric models, it is linear time cost. For kernel density estimation, the model learning cost can be up to quadratic.

3.3 Clustering-Based Methods

Possible scenarios:

- Does the instance belong to any cluster?
 - If no, then an outlier.
- Is there a large distance between the instance and the cluster to which it is closest?
 - If yes, then an outlier.
- Is the instance part of a small or sparse cluster?
 - If yes, then the whole cluster is outlying.

Usually **two steps**: (1) clustering, (2) outlier score measuring.

- Use existing outlier detection methods for the second step.

3.3 Clustering-Based Methods

A review of k-mean clustering:

- Given a set of observations $(\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_N)$, k-means clustering aims to partition the N observations into $k (\leq n)$ sets $S = \{S_1, S_2, \dots, S_k\}$.
- Formally, the objective is to find:

$$\operatorname{argmin}_S \sum_{i=1}^k \sum_{\mathbf{x} \in S_i} \|\mathbf{x} - \boldsymbol{\mu}_i\|^2, \quad (2)$$

where $\boldsymbol{\mu}_i$ is the mean of S_i .

- This is equivalent to:

$$\operatorname{argmin}_S \sum_{i=1}^k \frac{1}{2|S_i|} \sum_{\mathbf{x}, \mathbf{y} \in S_i} \|\mathbf{x} - \mathbf{y}\|^2, \quad (3)$$

i.e., minimizing the pairwise squared deviations of points in the same cluster.

3.3 Clustering-Based Methods

Equivalence between the two objectives:

$$\operatorname{argmin}_S \sum_{i=1}^k \sum_{\mathbf{x} \in S_i} \|\mathbf{x} - \boldsymbol{\mu}_i\|^2$$
$$\operatorname{argmin}_S \sum_{i=1}^k \frac{1}{2|S_i|} \sum_{\mathbf{x}, \mathbf{y} \in S_i} \|\mathbf{x} - \mathbf{y}\|^2$$

3.3 Clustering-Based Methods

A review of k-means clustering:

- Initiate the k mean vectors $\{\mu_1, \mu_2, \dots, \mu_k\}$.
- Alternate between the two steps as below.
 - 1 **Assignment:** Assign each instance to the cluster with the nearest mean vector:

$$S_i^{(t)} = \{\mathbf{x}_p : \|\mathbf{x}_p - \mu_i\| \leq \|\mathbf{x}_p - \mu_j\|, \forall j, 1 \leq j \leq k\} \quad (4)$$

- 2 **Update:** Recalculate mean vectors

$$\mathbf{u}_i = \frac{1}{|S_i^{(t)}|} \sum_{\mathbf{x}_p \in S_i^{(t)}} \mathbf{x}_p \quad (5)$$

- The algorithm converges when the assignments no longer change.

3.3 Clustering-Based Methods

Some additional details in k-means:

- How to initialize?
 - The Forge method.
 - The Random Partition method.
- How to find the best k ?
 - Through validation.
 - We could use the Gap Statistic ².
- Limitations of k-mean?

²"Estimating the number of clusters in a data set via the gap statistic."
Journal of the Royal Statistical Society. 2001.

3.3 Clustering-Based Methods

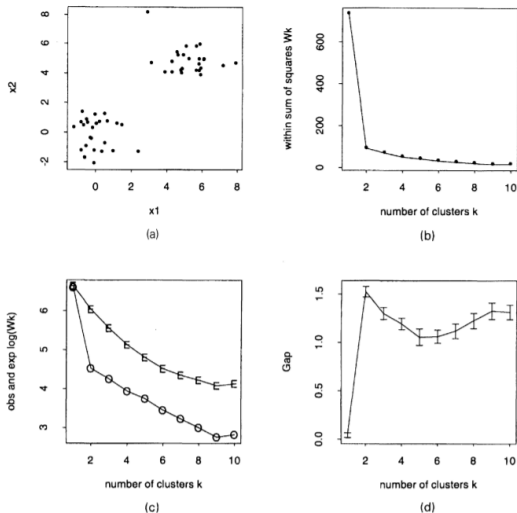


Fig. 1. Results for the two-cluster example: (a) data; (b) within sum of squares function W_k ; (c) functions $\log(W_k)$ (O) and $E_n^*(\log(W_k))$ (E); (d) gap curve

3.3 Clustering-Based Methods


Properties of clustering-based methods:

- Obtain a summary of data
- Good explainability
- The first step could be time consuming
- The detection performance largely depends on clustering quality
 - Gaussian Mixture Model
 - Hierarchical Clustering
 - Spectral Clustering
 - Deep Model Based Clustering
 - ...

3.5 iForest

Isolation Forest (iForest) ³

- **Isolation** means “separating an instance from the rest of the instances”.
- Easily being separated \rightarrow More likely to be an outlier.



³"Isolation forest." ICDM. 2008.


3.5 iForest

Algorithm 1 : $iForest(X, t, \psi)$


Inputs: X - input data, t - number of trees, ψ - subsampling size

Output: a set of t *iTrees*

- 1: **Initialize** $Forest$
 - 2: set height limit $l = \text{ceiling}(\log_2 \psi)$
 - 3: **for** $i = 1$ to t **do**
 - 4: $X' \leftarrow \text{sample}(X, \psi)$
 - 5: $Forest \leftarrow Forest \cup iTree(X', 0, l)$
 - 6: **end for**
 - 7: **return** $Forest$
-



(a) Original sample
(4096 instances)



(b) Sub-sample
(128 instances)

3.5 iForest

Algorithm 2 : $iTree(X, e, l)$

Inputs: X - input data, e - current tree height, l - height limit

Output: an iTree

```
1: if  $e \geq l$  or  $|X| \leq 1$  then
2:   return  $exNode\{Size \leftarrow |X|\}$ 
3: else
4:   let  $Q$  be a list of attributes in  $X$ 
5:   randomly select an attribute  $q \in Q$ 
6:   randomly select a split point  $p$  from  $max$  and  $min$ 
       values of attribute  $q$  in  $X$ 
7:    $X_l \leftarrow filter(X, q < p)$ 
8:    $X_r \leftarrow filter(X, q \geq p)$ 
9:   return  $inNode\{Left \leftarrow iTree(X_l, e + 1, l),$ 
10:                 $Right \leftarrow iTree(X_r, e + 1, l),$ 
11:                 $SplitAtt \leftarrow q,$ 
12:                 $SplitValue \leftarrow p\}$ 
13: end if
```

There are two types of nodes: $inNode$ and $exNode$.

Not always separate each point.

3.5 iForest

Complexity $O(nt\log\psi)$

Algorithm 3 : $PathLength(x, T, e)$

Inputs : x - an instance, T - an iTree, e - current path length;
to be initialized to zero when first called

Output: path length of x

```
1: if  $T$  is an external node then
2:   return  $e + c(T.size)$   $\{c(.)$  is defined in Equation 1 $\}$ 
3: end if
4:  $a \leftarrow T.splitAtt$ 
5: if  $x_a < T.splitValue$  then
6:   return  $PathLength(x, T.left, e + 1)$ 
7: else  $\{x_a \geq T.splitValue\}$ 
8:   return  $PathLength(x, T.right, e + 1)$ 
9: end if
```

3.5 iForest

The outlier score is an instance \mathbf{x} is:

$$s(\mathbf{x}, \psi) = 2^{-\frac{E(h(\mathbf{x}))}{c(\psi)}} \quad (6)$$

- $h(\mathbf{x})$: a single path length from root to \mathbf{x} .
- $c(\psi)$: normalization
- iTrees have an equivalent structure to Binary Search Tree (BST).
- The estimation of $E(h(\mathbf{x}))$ for external node terminations is the same as the unsuccessful search in BST.
- $c(\psi) =$ the average path length of unsuccessful search in BST.

3.5 iForest

The outlier score is an instance \mathbf{x} is:


$$s(\mathbf{x}, \psi) = 2^{-\frac{E(h(\mathbf{x}))}{c(\psi)}} \quad (7)$$

- If instances return $s \approx 1$, then they are outliers.
- If instances have $s \ll 0.5$, then they are quite safe to be seen as normal instances.
- If all the instances return $s \approx 0.5$, then the entire sample does not really have any distinct outliers.

3.5 iForest

Two variables to choose in iForest:

- number of trees to build t (default: 100)
- sub-sampling size ψ (default: 256)



4 Outlier Detection in High-Dimensional Data

Dimensionality Curse: Various phenomena that arise when analyzing data in high-dimensional spaces that do not occur in low-dimensional ones ⁴.

- The *pairwise distances* between different samples in the space converging to the same value as the dimensionality of the data increases.
- As the dimensionality increases, the distance between objects may be heavily dominated by noise.
- The *proximity* or *similarity* between samples may not be qualitatively relevant in higher dimensions.

⁴https://en.wikipedia.org/wiki/Curse_of_dimensionality

4 Outlier Detection in High-Dimensional Data

Dimensionality Reduction: compress data

- Classical approaches
 - E.g., Principle Component Analysis (PCA)
- Deep learning approaches.
 - E.g., Robust Deep Autoencoder (RDA) ⁵

⁵Zhou, Chong, and Randy C. Paffenroth. "Anomaly detection with robust deep autoencoders." KDD. 2017.

4.1 PCA

- A feature/dimensionality reduction technique.
- A relevant (but different) technique is feature selection.
- The goal of PCA is to transform the high-dimensional features to a lower-dimensional space.
- PCA is unsupervised.

What is a good data transformation?

4.1 PCA

What is a good data compression (transformation)?

- High compression rate (smaller data size).
- Key characteristics are maintained (keep more information).

In the deep learning era, the transformed/compressed data is also called as:

- the “representation” of data
- “latent representation”
- “embeddings”

4.1 PCA

- Let us start from the simple case where PCA projects the features to a **1-dimensional space**.
- Suppose we have n p -dimensional ($p > 1$) instances, $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n \in \mathbb{R}^p$.
- We want to learn a mapper $\mathbf{a} \in \mathbb{R}^p$ that $\mathbf{a}^\top \mathbf{x}_i = z_i$. Here z_i is where \mathbf{x}_i locates in the 1-dimensional space.
- PCA tries to solve:

$$\mathbf{a}^* = \underset{\|\mathbf{a}\|=1}{\operatorname{argmax}} \underbrace{\frac{1}{n} (z_i - \bar{z})^2}_{\text{variance}}, \quad (8)$$

where PCA tries to find the projection with the maximum variance in reduced data.

4.1 PCA

- PCA tries to solve:

$$\mathbf{a}^* = \underset{\|\mathbf{a}\|=1}{\operatorname{argmax}} \underbrace{\frac{1}{n} (z_i - \bar{z})^2}_{\text{variance}}. \quad (9)$$

- Q1: Why $\|\mathbf{a}\| = 1$?
- Q2: What does the objective imply?
 - Trying to make the data points distinguishable in the new space.
 - Why?

4.1 PCA

Since

$$\bar{z} = \frac{1}{n} \sum_{i=1}^n z_i = \frac{1}{n} \sum_{i=1}^n \mathbf{a}^T \mathbf{x}_i = \mathbf{a}^T \left(\frac{1}{n} \sum_{i=1}^n \mathbf{x}_i \right) = \mathbf{a}^T \bar{\mathbf{x}},$$

the problem can be written as

$$\begin{aligned} \mathbf{a}^* &= \arg \max_{\|\mathbf{a}\|=1} \frac{1}{n} \sum_{i=1}^n (z_i - \bar{z})^2 \\ &= \arg \max_{\|\mathbf{a}\|=1} \frac{1}{n} \sum_{i=1}^n (\mathbf{a}^T \mathbf{x}_i - \bar{z})^2 \\ &= \arg \max_{\|\mathbf{a}\|=1} \frac{1}{n} \sum_{i=1}^n (\mathbf{a}^T \mathbf{x}_i - \mathbf{a}^T \bar{\mathbf{x}})^2 \\ &= \arg \max_{\|\mathbf{a}\|=1} \frac{1}{n} \sum_{i=1}^n \mathbf{a}^T (\mathbf{x}_i - \bar{\mathbf{x}}) (\mathbf{x}_i - \bar{\mathbf{x}})^T \mathbf{a} \\ &= \arg \max_{\|\mathbf{a}\|=1} \mathbf{a}^T \underbrace{\left(\frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}}) (\mathbf{x}_i - \bar{\mathbf{x}})^T \right)}_{p \times p \text{ covariance matrix}} \mathbf{a} \\ &= \arg \max_{\|\mathbf{a}\|=1} \mathbf{a}^T \mathbf{C} \mathbf{a}, \end{aligned}$$

where $\mathbf{C} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}}) (\mathbf{x}_i - \bar{\mathbf{x}})^T$, denotes the covariance matrix.

4.1 PCA

The problem is

$$\begin{aligned} & \max_{\mathbf{a}} \mathbf{a}^\top \mathbf{C} \mathbf{a} \\ & s.t., \|\mathbf{a}\|^2 = 1. \end{aligned} \tag{10}$$

How to solve this?

4.1 PCA

Eigen-Decomposition

A square $n \times n$ matrix \mathbf{C} with n linearly independent eigenvectors can be factorized as:

$$\mathbf{C} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^{-1} \quad (11)$$

- \mathbf{U} is a $n \times n$ matrix whose columns are eigenvectors of \mathbf{C} ,
- $\mathbf{\Lambda}$ is a diagonal matrix whose diagonal elements are the corresponding eigenvalues.

If \mathbf{C} is symmetric, then

$$\mathbf{C} = \mathbf{U}\mathbf{\Lambda}\mathbf{U}^T = \sum_{i=1}^n \lambda_i \mathbf{u}_i \mathbf{u}_i^T. \quad (12)$$

4.1 PCA

Q: What about mapping data into a k -dimensional space?

A: The PCA problem is then defined as:

$$\mathbf{A}^* = \operatorname{argmax}_{\mathbf{A} \in \mathbb{R}^{p \times k}: \mathbf{A}^\top \mathbf{A} = \mathbf{I}_k} \operatorname{trace}(\mathbf{A}^\top \mathbf{C} \mathbf{A}) \quad (13)$$

where $\mathbf{A} = [\mathbf{a}_1, \mathbf{a}_2, \dots, \mathbf{a}_k] \in \mathbb{R}^{p \times k}$. PCA requires different projection vectors to be orthogonal.

$$\operatorname{trace}(\mathbf{A}^\top \mathbf{C} \mathbf{A}) = \sum_{i=1}^k (\mathbf{a}_i^\top \mathbf{C} \mathbf{a}_i), \quad (14)$$

which is a natural extension of 1-D PCA.

Q: How to solve this?

4.1 PCA

When $k = 2$, we want to solve

$$\begin{aligned} & \max_{\mathbf{a}_2} \mathbf{a}_2^T \mathbf{C} \mathbf{a}_2 \\ & s.t., \|\mathbf{a}_i\|^2 = 1 \\ & \mathbf{a}_2^T \mathbf{a}_1 = 0. \end{aligned} \tag{15}$$

How to solve this?

4.1 PCA

We then have the following theorem.

Theorem. (*Ky Fan*) Let $\mathbf{H} \in \mathbb{R}^{n \times n}$ be a symmetric matrix with eigenvalues

$$\lambda_1 \geq \lambda_2 \geq \cdots \geq \lambda_n,$$

and the corresponding eigenvectors $\mathbf{U} = [\mathbf{u}_1, \dots, \mathbf{u}_n]$. Then

$$\lambda_1 + \cdots + \lambda_k = \max_{\mathbf{A} \in \mathbb{R}^{n \times k}: \mathbf{A}^T \mathbf{A} = \mathbf{I}_k} \text{trace}(\mathbf{A}^T \mathbf{H} \mathbf{A}).$$

And the optimal \mathbf{A}^* is given by $\mathbf{A}^* = [\mathbf{u}_1, \dots, \mathbf{u}_k] \mathbf{Q}$ with \mathbf{Q} an arbitrary orthogonal matrix. ■

4.1 PCA

- ① Note that in Eqn. (2), the covariance matrix \mathbf{C} is a symmetric matrix. Given the above theorem, we directly obtain

$$\begin{aligned}\lambda_1 + \cdots \lambda_k &= \arg \max_{\mathbf{A} \in \mathbb{R}^{n \times k}: \mathbf{A}^T \mathbf{A} = \mathbf{I}_k} \text{trace} \left(\mathbf{A}^T \mathbf{C} \mathbf{A} \right), \\ \mathbf{A}^* &= [\mathbf{u}_1, \dots, \mathbf{u}_k] \mathbf{Q},\end{aligned}$$

where $\lambda_1, \dots, \lambda_k$ are the k largest eigenvalues of the covariance matrix \mathbf{C} , and the solution \mathbf{A}^* is the matrix whose columns are corresponding eigenvectors.

- ② It also follows from the above theorem that solutions to PCA are not unique, and they differ by an orthogonal matrix. We used the special case where $\mathbf{Q} = \mathbf{I}$, i.e., $\mathbf{A}^* = [\mathbf{u}_1, \dots, \mathbf{u}_k]$.

4.1 PCA for Outlier Detection

Using PCA, we obtain the mappers $\mathbf{A} = \{\mathbf{u}_i\}$, $1 \leq i \leq k$. An original data instance \mathbf{x} will be mapped to $\mathbf{z} = [z_1, z_2, \dots, z_k]$. After mapping all the data, outlier detection could be done in the \mathbf{z} -space.

- Apply the detection methods we already discussed.
- Choose the mappers corresponding to higher or lower eigenvalues?
- How many k 's to choose?

4.1 PCA for Outlier Detection

There is another way to understand PCA from the **data reconstruction** perspective.

- Assume \mathbf{X} is centered.
- Data mapping: $\mathbf{Z} = \mathbf{A}^T \mathbf{X}$.
- Reconstructed data: $\tilde{\mathbf{X}} = \mathbf{A} \mathbf{Z} = \mathbf{A} \mathbf{A}^T \mathbf{X}$.
- Reconstruction error: $\|\mathbf{X} - \tilde{\mathbf{X}}\|_F = \|\mathbf{X} - \mathbf{A} \mathbf{A}^T \mathbf{X}\|_F$.
- It has also been shown that:

$$\mathbf{A} \mathbf{A}^T \mathbf{X} = \mathbf{B}^* = \underset{\mathbf{B}: \text{rank}(\mathbf{B}) \leq k}{\text{argmin}} \|\mathbf{X} - \mathbf{B}\|_F. \quad (16)$$

PCA projects the high-dimensional features to a lower-dimensional space with minimum reconstruction error.

- Keep most of the information = Be able to recover.
- Instances with large reconstruction errors $\|\mathbf{x} - \mathbf{A} \mathbf{A}^T \mathbf{x}\|_2$ are outliers.

4.1 PCA for Outlier Detection

In previous parts, we need to compute \mathbf{C} first before doing PCA. This step could be costly if p or n is large. Can this step be avoided?

Singular Value Decomposition:

The singular value decomposition (SVD) of an $m \times n$ real matrix \mathbf{R} (we assume $m \geq n$) can be written as:

$$\mathbf{R} = \mathbf{U}\mathbf{\Sigma}\mathbf{V}^T, \quad (17)$$

where $\mathbf{U} \in \mathbb{R}^{m \times m}$, $\mathbf{\Sigma} \in \mathbb{R}^{m \times n}$, and $\mathbf{V} \in \mathbb{R}^{n \times n}$, $\mathbf{U}\mathbf{U}^T = \mathbf{U}^T\mathbf{U} = \mathbf{I}_{m \times m}$, and $\mathbf{V}\mathbf{V}^T = \mathbf{V}^T\mathbf{V} = \mathbf{I}_{n \times n}$.

4.1 PCA for Outlier Detection

1. Pre-process data.

$$\mathbf{X} = [\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n] \in \mathbb{R}^{p \times n},$$

$$\bar{\mathbf{X}} = \frac{1}{n} \mathbf{X} \mathbf{1}_n \in \mathbb{R}^{p \times 1},$$

$$\tilde{\mathbf{X}} = (\mathbf{X} - \bar{\mathbf{X}} \mathbf{1}_n^T) \in \mathbb{R}^{p \times n},$$

2. We know $\mathbf{C} = \frac{1}{n} \sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})(\mathbf{x}_i - \bar{\mathbf{x}})^T = \frac{1}{n} \tilde{\mathbf{X}} \tilde{\mathbf{X}}^T$.
3. If the SVD of $\tilde{\mathbf{X}}$ is $\tilde{\mathbf{X}} = \mathbf{U} \tilde{\Sigma} \mathbf{V}$, the columns of \mathbf{U} (left-singular vectors) are eigenvectors of $\tilde{\mathbf{X}} \tilde{\mathbf{X}}^T$.
4. We do not perform eigen-value decomposition of \mathbf{C} , but do singular-value decomposition of $\tilde{\mathbf{X}}$, to obtain the mappers \mathbf{U} .

4.1 PCA for Outlier Detection

A more efficient way of computing PCA can be described as:

$$\underbrace{\mathbf{X}}_{p \times n} \rightarrow \underbrace{\tilde{\mathbf{X}}}_{p \times n} \rightarrow \tilde{\mathbf{X}} = \mathbf{U} \tilde{\Sigma} \mathbf{V}^T \rightarrow \mathbf{G} = \mathbf{U}_k \in \mathbb{R}^{p \times k}$$

$$\mathbf{G}^T \tilde{\mathbf{X}} = \mathbf{Z} \in \mathbb{R}^{k \times n}$$

4.1 Connection between PCA and Autoencoders




Figure 3: An autoendoer.

- What does each module in PCA correspond to?
- How to go deeper?

4.2 Outlier Detection with Deep Autoencoders

Deep Autoencoders

- A feed-forward multi-layer neural network.
- The desired output is the input itself.
 - Looks trivial?
 - Making hidden representation to be low-dimensional.
- The loss function of training an autoencoder:

$$\min_{D,E} \|\mathbf{X} - D(E(\mathbf{X}))\| \quad (18)$$

4.2 Outlier Detection with Deep Autoencoders

Robust PCA (RPCA)

- Try to reduce the sensitivity of PCA to outliers. Intuition?
- Split the data matrix into two parts

$$\mathbf{X} = \mathbf{L} + \mathbf{S}, \quad (19)$$

where \mathbf{L} is a low-rank matrix, and \mathbf{S} is a sparse matrix.

4.2 Outlier Detection with Deep Autoencoders

Robust PCA (RPCA)

- The problem can be formulated as:

$$\begin{aligned} \min_{\mathbf{L}, \mathbf{S}} \quad & \|\mathbf{L}\|_* + \lambda \|\mathbf{S}\|_1 \\ \text{s.t.} \quad & \|\mathbf{X} - \mathbf{L} - \mathbf{S}\|_F^2 = 0, \end{aligned} \tag{20}$$

where $\|\cdot\|_*$ is the nuclear norm, $\|\cdot\|_1$ is the L1 norm.

- If we let D and E to be linear mappers, then PCA is the optimal solution for recovering \mathbf{L} .

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

- The problem can be formulated as:

$$\begin{aligned} \min_{\mathbf{L}, \mathbf{S}} \quad & \|\mathbf{L} - D_{\theta}(E_{\theta}(\mathbf{L}))\|_2 + \lambda \|\mathbf{S}\|_1 \\ \text{s.t.} \quad & \mathbf{X} - \mathbf{L} - \mathbf{S} = 0, \end{aligned} \tag{21}$$

where both the encoder and the decoder are parameterized by θ .

- The first term encourages \mathbf{L} to contain the “main features” in data.
- λ plays an essential role:
 - a smaller $\lambda \rightarrow$ encourage more data into \mathbf{S}
 - a larger $\lambda \rightarrow$ include less data in \mathbf{S}

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

- How to choose D_θ and E_θ ?
 - Any commonly used deep autoencoder architectures could be used.
 - The paper choose a simple architecture:

$$\begin{aligned}E_\theta(\mathbf{x}) &= g(\mathbf{W}\mathbf{x} + b_E) \\ D_\theta(\mathbf{x}) &= g(\mathbf{W}^\top E_\theta(\mathbf{x}) + b_D)\end{aligned}\tag{22}$$

- $g(x) = \frac{1}{1+\exp(-x)}$, a nonlinear transformation.

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

- Is $\|S\|_1$ always a good choice?
- What if each anomaly point is not a matrix entry, but the whole row or column?




Figure 4: Two examples of group anomalies.

4.2 Outlier Detection with Deep Autoencoders

Regularization

- $\min \|x\|_1$ vs $\min \|x\|_2$
- Which one encourages a sparse x ?
- An example of taxes.
- Can we combine them?

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

- $l_{2,1}$ normalization encourages structured/group sparsity.
- $\|\mathbf{X}\|_{2,1} = \sum_{j=1}^n \|\mathbf{x}_j\|_2 = \sum_{j=1}^n (\sum_{i=1}^m |x_{i,j}|^2)^{1/2}$.
- $\|\mathbf{X}^\top\|_{2,1}$?




Figure 5: Two examples of group anomalies.

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

- The problem can be formulated as:

$$\begin{aligned} \min_{\mathbf{L}, \mathbf{S}} \quad & \|\mathbf{L} - D_{\theta}(E_{\theta}(\mathbf{L}))\|_2 + \lambda(\|\mathbf{S}\|_1 \text{ or } \|\mathbf{S}\|_{2,1} \text{ or } \|\mathbf{S}^T\|_{2,1}) \\ \text{s.t.} \quad & \mathbf{X} - \mathbf{L} - \mathbf{S} = 0, \end{aligned}$$

- Just denoising: choose $\|\mathbf{S}\|_1$.
- To find abnormal data instances: choose $\|\mathbf{S}^T\|_{2,1}$.

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

- How to solve the problem?

$$\min_{\mathbf{L}, \mathbf{S}} \|\mathbf{L} - D_{\theta}(E_{\theta}(\mathbf{L}))\|_2 + \lambda(\|\mathbf{S}\|_1 \text{ or } \|\mathbf{S}\|_{2,1} \text{ or } \|\mathbf{S}^T\|_{2,1})$$
$$s.t. \mathbf{X} - \mathbf{L} - \mathbf{S} = 0,$$

- Alternating Direction Method of Multipliers (ADMM):
 - Divide the objective into several pieces.
 - Optimizes one of the pieces while keeping the others fixed.
- How to update \mathbf{L} with respect to $\min_{\mathbf{L}} \|\mathbf{L} - D_{\theta}(E_{\theta}(\mathbf{L}))\|_2$?
- How to update \mathbf{S} with respect to $\min_{\mathbf{S}} \|\mathbf{S}\|_1$?

4.2 Outlier Detection with Deep Autoencoders

Robust Deep Autoencoders (RDA)

The algorithm of solving

$$\min_{\mathbf{L}, \mathbf{S}} \|\mathbf{L} - D_{\theta}(E_{\theta}(\mathbf{L}))\|_2 + \lambda(\|\mathbf{S}\|_1 \text{ or } \|\mathbf{S}\|_{2,1} \text{ or } \|\mathbf{S}^T\|_{2,1})$$

s.t. $\mathbf{X} - \mathbf{L} - \mathbf{S} = 0.$

- ① Initialize θ, \mathbf{S}
- ② while not converge:
 - ① $\mathbf{L} = \mathbf{X} - \mathbf{S}$
 - ② Minimize $\|\mathbf{L} - D_{\theta}(E_{\theta}(\mathbf{L}))\|_2$ with gradient descent
 - ③ $\mathbf{L} \leftarrow D_{\theta}(E_{\theta}(\mathbf{L}))$
 - ④ $\mathbf{S} = \mathbf{X} - \mathbf{L}$
 - ⑤ Optimize \mathbf{S} using a proximal operator
 - ⑥ Check convergence.

4.2 Outlier Detection with Deep Autoencoders

Experiment Design

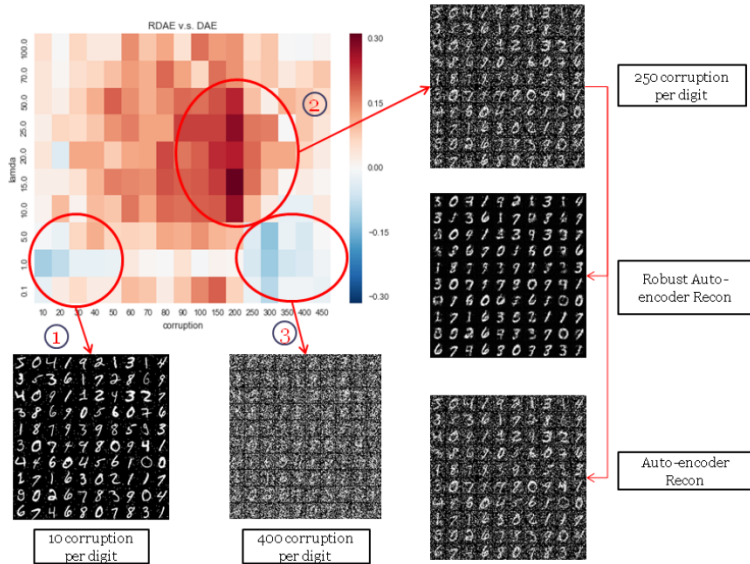
- Datasets
- Research Problems
 - Baseline methods
 - Metrics
 - Visualization
- Effect of hyperparameters or modules

4.2 Outlier Detection with Deep Autoencoders

Experiment Design in RDA

- Datasets: MNIST
- Research Problem 1: Whether excluding from data an independent sparse matrix using $\|\mathbf{S}\|_1$ could help the autoencoder in capturing the essential patterns in data.
 - Baseline methods: Autoencoder
 - Metrics: Classification error
- Experiment design:
 - Extract hidden features from $E_\theta(\mathbf{L})$.
 - Use the hidden features to build a classifier to classify digits.
 - Lower error \rightarrow Better model.

4.2 Outlier Detection with Deep Autoencoders




4.2 Outlier Detection with Deep Autoencoders

Experiment Design in RDA

- Datasets: MNIST
- Research Problem 2: Is $\|\mathbf{S}\|_{2,1}$ regularized RDA effective in detecting outliers?
 - Baseline method: Isolation Forest (IS)
 - Evaluation metric: F1-score
- Experiment design:
 - Compare detected outliers with the ground truths.
 - 0.64 (RDA) vs 0.37 (IS)


4.2 Outlier Detection with Deep Autoencoders



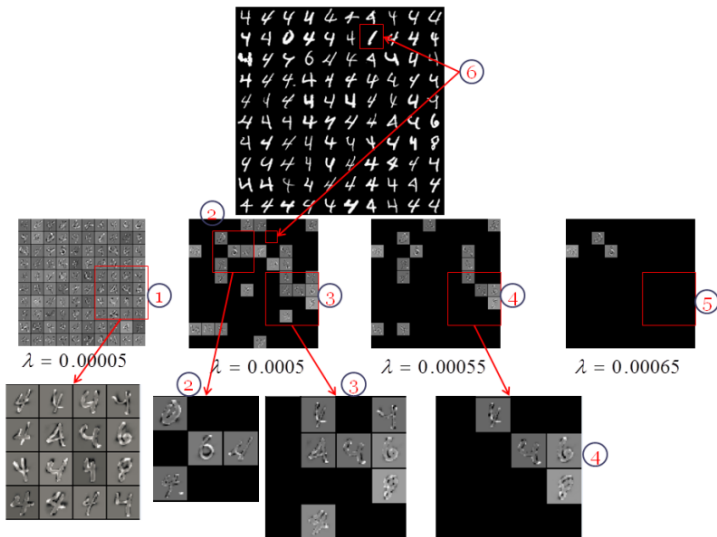
4.2 Outlier Detection with Deep Autoencoders

Experiment Design in RDA

- Datasets: MNIST
- Research Problem 3: Effect of λ ?



4.2 Outlier Detection with Deep Autoencoders



4.2 Outlier Detection with Deep Autoencoders